

Reinforcement Learning

DQN & SARSA

Naeemullah Khan

naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

June 30, 2025

1. Learning Outcomes
2. Motivation DQN
3. Deep Q-Network (DQN) Overview
4. Training the Q-network
5. DQN Algorithm
6. Double DQN
7. On-Policy vs Off-Policy Learning
8. SARSA
9. Case Study: Playing Atari Games
10. Summary
11. References

By the end of this session, you will be able to:

- ▶ Understand and implement **Deep Q-Networks (DQN)**.
- ▶ Grasp the **SARSA algorithm** and how it contrasts with Q-learning.
- ▶ Analyze the **differences between off-policy and on-policy learning**.
- ▶ Evaluate the strengths and weaknesses of both algorithms.
- ▶ Understand practical applications and limitations of DQN and SARSA.

- ▶ **Tabular Q-learning does not scale** to large or continuous state spaces.
- ▶ **Function approximation** is needed to generalize across similar states; deep learning provides powerful tools for this.
- ▶ **Deep Q-Networks (DQN)** combine Q-learning with deep neural networks, enabling reinforcement learning for high-dimensional inputs such as images.
- ▶ **Breakthrough:** DQN achieved human-level performance on Atari 2600 games directly from raw pixel inputs.

Reinforcement Learning: **Deep Q-Network (DQN) Overview**

- ▶ So far, we've been assuming a tabular representation of Q : one entry for every state/action pair
- ▶ What's the problem with this?

- ▶ So far, we've been assuming a tabular representation of Q : one entry for every state/action pair
- ▶ What's the problem with this?
- ▶ Not scalable. Must compute $Q(s, a)$ for every state-action pair. If state is e.g. current game state pixels, computationally infeasible to compute for entire state space!

- ▶ So far, we've been assuming a tabular representation of Q : one entry for every state/action pair
- ▶ What's the problem with this?
- ▶ Not scalable. Must compute $Q(s, a)$ for every state-action pair. If state is e.g. current game state pixels, computationally infeasible to compute for entire state space!
- ▶ **Solution:** use a function approximator to estimate $Q(s, a)$. E.g. a neural network!

$$Q(s, a; \theta) \approx Q^*(s, a)$$

- ▶ So far, we've been assuming a tabular representation of Q : one entry for every state/action pair
- ▶ What's the problem with this?
- ▶ Not scalable. Must compute $Q(s, a)$ for every state-action pair. If state is e.g. current game state pixels, computationally infeasible to compute for entire state space!
- ▶ **Solution:** use a function approximator to estimate $Q(s, a)$. E.g. a neural network!

$$Q(s, a; \theta) \approx Q^*(s, a)$$

- ▶ If the function approximator is a deep neural network \Rightarrow Deep Q-Learning (DQN)

- ▶ Remember: want to find a Q-function that satisfies the Bellman Equation:

$$Q^*(s, a) = \mathbb{E}_{p(s'|s,a)} \left[r(s, a) + \gamma \max_{a'} Q^*(s', a') | s, a \right]$$

- ▶ Forward Pass:

$$\mathcal{L}(\theta) = \mathbb{E}_{s,a \sim p(\cdot)} [(y - Q(s, a; \theta))^2]$$

$$\text{where } y = \mathbb{E}_{p(s'|s,a)} \left[r(s, a) + \gamma \max_{a'} Q(s', a'; \theta) | s, a \right]$$

- ▶ Backward Pass:

- Gradient update (with respect to Q-function parameters θ):

-2cm-2cm

$$\nabla_{\theta} \mathcal{L}(\theta) = \mathbb{E}_{s,a \sim p(\cdot), p(s'|s,a)} \left[(r(s, a) + \gamma \max_{a'} Q(s', a'; \theta) - Q(s, a; \theta)) \nabla_{\theta} Q(s, a; \theta) \right]$$

Reinforcement Learning: **Training the Q-network**

- ▶ Learning from batches of consecutive samples is problematic:
 - Samples are correlated \Rightarrow inefficient learning.
 - The current Q-network parameters determine the next training samples (e.g., if the maximizing action is to move left, training samples will be dominated by transitions from the left-hand side), which can lead to bad feedback loops.

- ▶ Learning from batches of consecutive samples is problematic:
 - Samples are correlated \Rightarrow inefficient learning.
 - The current Q-network parameters determine the next training samples (e.g., if the maximizing action is to move left, training samples will be dominated by transitions from the left-hand side), which can lead to bad feedback loops.
- ▶ These problems can be addressed using experience replay:
 - Maintain a replay memory buffer of transitions (s_t, a_t, r_t, s_{t+1}) as episodes are played.
 - Train the Q-network on random minibatches of transitions sampled from the replay memory, instead of using consecutive samples.

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for

- ▶ DQN, as stated above, will have learning problems because the target and the prediction are not independent as they both rely on the same network.
- ▶ It's like a dog chasing it's own tail.

- ▶ DQN, as stated above, will have learning problems because the target and the prediction are not independent as they both rely on the same network.
- ▶ It's like a dog chasing its own tail.
- ▶ **Solution:** Use two separate Q-value estimators, each of which is used to update the other.
- ▶ The target values are calculated using a target Q-network. The target Q-network's parameters are updated to the current networks every C time steps.
- ▶ Target network prevents the network from spiraling around.

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

With probability ε select a random action a_t

otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action a_t in emulator and observe reward r_t and image x_{t+1}

Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

Every C steps reset $\hat{Q} = Q$

End For

End For

⁰Minh et al. <https://www.nature.com/articles/nature14236>

Reinforcement Learning: **Double DQN**

- ▶ Q-learning suffers from a maximization bias.
- ▶ This is because the update target is $r + \max_a Q^*(s, a)$. If Q-value is slightly overestimated then this error gets compounded.

- ▶ Q-learning suffers from a maximization bias.
- ▶ This is because the update target is $r + \max_a Q^*(s, a)$. If Q-value is slightly overestimated then this error gets compounded.
- ▶ **Solution:** Decompose the max operation in the target into action selection and action evaluation.
- ▶ Use the current network to select the max action for the next state and then use the target network to get the target Q-value for that action.
- ▶ Using these independent estimators, we can have unbiased Q-value estimates of the actions selected using the opposite estimator.
- ▶ We can thus avoid maximization bias by disentangling our updates from biased estimates.

Algorithm 1: Double DQN Algorithm.

input : \mathcal{D} – empty replay buffer; θ – initial network parameters, θ^- – copy of θ
input : N_f – replay buffer maximum size; N_b – training batch size; N^- – target network replacement freq.
for episode $e \in \{1, 2, \dots, M\}$ **do**
 Initialize frame sequence $\mathbf{x} \leftarrow ()$
 for $t \in \{0, 1, \dots\}$ **do**
 Set state $s \leftarrow \mathbf{x}$, sample action $a \sim \pi_B$
 Sample next frame x^t from environment \mathcal{E} given (s, a) and receive reward r , and append x^t to \mathbf{x}
 if $|\mathbf{x}| > N_f$ **then** delete oldest frame $x_{t_{\min}}$ from \mathbf{x} **end**
 Set $s' \leftarrow \mathbf{x}$, and add transition tuple (s, a, r, s') to \mathcal{D} ,
 replacing the oldest tuple **if** $|\mathcal{D}| \geq N_f$
 Sample a minibatch of N_b tuples $(s, a, r, s') \sim \text{Unif}(\mathcal{D})$
 Construct target values, one for each of the N_b tuples:
 Define $a^{\max}(s'; \theta) = \arg \max_a Q(s', a'; \theta)$

$$y_j = \begin{cases} r & \text{if } s' \text{ is terminal} \\ r + \gamma Q(s', a^{\max}(s'; \theta); \theta^-), & \text{otherwise.} \end{cases}$$
 Do a gradient descent step with loss $\|y_j - Q(s, a; \theta)\|^2$
 Replace target parameters $\theta^- \leftarrow \theta$ every N^- steps
 end
end

⁰<https://leejung.github.io/posts/Dueling-DQN/>

► On-Policy:

- Agent learns by doing.
- Follows one policy for both acting and learning.
- Example: Tries out actions, learns from its own experience.

► Off-Policy:

- Agent learns from others or past experiences.
- Can follow one policy but learn about another.
- Example: Watches someone else, learns what would have happened if it acted differently.

Reinforcement Learning: **SARSA**

- **SARSA** stands for:

State
→ Action
→ Reward
→ State'
→ Action'

- **On-policy learning:** Learns Q-values by following the current behavior policy.
- **Update Rule:**

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]$$

- ▶ Uses the action the agent actually takes, not just the best possible one.
- ▶ Learns from the agent's real experience, not hypothetical alternatives.
- ▶ Tends to be safer—useful in situations where mistakes are costly.
- ▶ Helps the agent improve its actual behavior step by step.

- ▶ The SARSA algorithm is a slight variation of the Q-Learning algorithm.
- ▶ Q-Learning is an **off-policy** method and uses a greedy approach to learn the Q-values.
- ▶ SARSA, on the other hand, is an **on-policy** method and uses the action performed by the current policy to update the Q-values.

- ▶ The SARSA algorithm is a slight variation of the Q-Learning algorithm.
- ▶ Q-Learning is an **off-policy** method and uses a greedy approach to learn the Q-values.
- ▶ SARSA, on the other hand, is an **on-policy** method and uses the action performed by the current policy to update the Q-values.

Q-Learning:
$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r(s_t, a_t) + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right]$$

- ▶ The SARSA algorithm is a slight variation of the Q-Learning algorithm.
- ▶ Q-Learning is an **off-policy** method and uses a greedy approach to learn the Q-values.
- ▶ SARSA, on the other hand, is an **on-policy** method and uses the action performed by the current policy to update the Q-values.

Q-Learning:
$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r(s_t, a_t) + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right]$$

SARSA:
$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r(s_t, a_t) + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t) \right]$$

- ▶ The update equation for SARSA depends on the current state, current action, reward obtained, next state, and next action.
- ▶ This observation led to the naming of the learning technique as **SARSA**, which stands for **State-Action-Reward-State-Action**, symbolizing the tuple (s, a, r, s', a') .

- ▶ The update equation for SARSA depends on the current state, current action, reward obtained, next state, and next action.
- ▶ This observation led to the naming of the learning technique as **SARSA**, which stands for **State-Action-Reward-State-Action**, symbolizing the tuple (s, a, r, s', a') .
- ▶ Similar to DQN, there is also a Deep SARSA variant.

Input : States, $s \in S$, Actions $a \in A(s)$, Initialize $Q(s, a)$, α , γ , π to an arbitrary policy (non-greedy)

Output: Optimal action value $Q(s, a)$ for each state-action pair

while *True* do

 for ($i = 0$; $i \leq \# \text{ of episodes}$; $i++$) do

 Initialize s

 Choose a from s , using policy derived from Q

 Repeat(for each step of episodes):

 Take action a ; observe reward, r , and next state, s'

 Choose action a' from state s' using policy derived from Q

$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$

$s \leftarrow s'$; $a \leftarrow a'$;

 until s is terminal

 end

end

⁰Lockery & Peters, *Adaptive learning by a target-tracking system*

Q-Learning vs SARSA: Comparison Table

Feature	Q-Learning	SARSA
Policy Type	Off-policy	On-policy
Target	$\max Q(s', a')$	$Q(s', a')$ actually taken
Exploration Aware	No	Yes
Risk Sensitivity	More aggressive	More conservative
Use Case	Goal-seeking behavior	Risk-aware behavior


Case Study: **Playing Atari Games**



- ▶ **Objective:** Complete the game with the highest score
- ▶ **State:** Raw pixel inputs of the game state
- ▶ **Action:** Game controls e.g. Left, Right, Up, Down
- ▶ **Reward:** Score increase/decrease at each time step


⁰[Mnih et al. NIPS Workshop 2013; Nature 2015]

$Q(s, a; \theta)$:
neural network
with weights θ



Last FC layer has 4-d
output (if 4 actions),
corresponding to $Q(s_t, a_1)$, $Q(s_t, a_2)$, $Q(s_t, a_3)$,
 $Q(s_t, a_4)$


Number of actions between 4-18
depending on Atari game




Current state s_t : 84x84x4 stack of last 4 frames
(after RGB->grayscale conversion, downsampling, and cropping)

⁰[Mnih et al. NIPS Workshop 2013; Nature 2015]

$Q(s, a; \theta)$:
neural network
with weights θ



← Last FC layer has 4-d output (if 4 actions), corresponding to $Q(s_t, a_1)$, $Q(s_t, a_2)$, $Q(s_t, a_3)$, $Q(s_t, a_4)$



Number of actions between 4-18 depending on Atari game

Current state s_t : 84x84x4 stack of last 4 frames
(after RGB->grayscale conversion, downsampling, and cropping)

<https://www.youtube.com/watch?v=V1eYniJ0Rnk>

⁰[Mnih et al. NIPS Workshop 2013; Nature 2015]

DDQN & SARSA: **Summary**

When to Use Which Algorithm?

Scenario	Recommended Algorithm
Environment is stochastic	SARSA (conservative)
Maximizing reward is the priority	Q-Learning or DQN
Large state space (e.g., images)	DQN
Limited computational resources	SARSA (simpler model)

- ▶ **Dueling DQN:** Separates value and advantage estimation.
- ▶ **Double DQN:** Reduces overestimation bias.
- ▶ **Prioritized Replay:** Samples important experiences more frequently.
- ▶ **Distributional RL:** Models the full return distribution, not just the expectation.
- ▶ **Safe RL:** Ensures safety in high-stakes environments (e.g., medical, robotics).

- ▶ A Markov Decision Process (MDP) is the mathematical formulation of the reinforcement learning problem, defined by $(\mathcal{S}, \mathcal{A}, \mathcal{R}, \mathbb{P}, \gamma)$.
- ▶ Each state satisfies the Markov property, i.e., the future is independent of the past given the present.
- ▶ The agent and the environment interact in a sequential loop. The policy π determines how the agent chooses actions.
- ▶ The value function estimates how good a state is, while the Q-value function estimates the quality of a state-action pair.
- ▶ The Bellman equation is a recursive formula for the Q-value function.

- ▶ Q-learning is an algorithm that repeatedly adjusts Q-values to minimize the Bellman error.
- ▶ When the Q-value function approximator is a deep neural network, we obtain Deep Q-Learning.
- ▶ DQN is powerful for complex, high-dimensional inputs but sensitive and data-hungry.
- ▶ SARSA is an on-policy variation of Q-learning.
- ▶ SARSA offers a safer, on-policy alternative, better suited for uncertain environments.
- ▶ Choice depends on task risk, dimensionality, and training stability.

DDQN & SARSA: **References**

- [1] Mnih, V., Kavukcuoglu, K., Silver, D., et al. (2015).
Human-level control through deep reinforcement learning.
Nature.
- [2] Van Hasselt, H., Guez, A., Silver, D. (2016).
Deep Reinforcement Learning with Double Q-learning.
AAAI.
- [3] Wang, Z., Schaul, T., Hessel, M., et al. (2016).
Dueling Network Architectures for Deep RL.
ICML.

- [4] Rummery, G. A., & Niranjan, M. (1994).
On-line Q-learning using connectionist systems.
Technical Report.
- [5] Sutton, R. S., & Barto, A. G. (2018).
Reinforcement Learning: An Introduction.
- [6] Berkeley CS285 Deep RL Lectures.
<https://rail.eecs.berkeley.edu/deeprlcourse/>
- [7] David Silver's RL Series.
<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>
- [8] Emma Brunskill, Stanford CS234: Reinforcement Learning.
<https://web.stanford.edu/class/cs234/>

- [9] Fei-Fei Li, Yunzhu Li & Ruohan Gao, Stanford CS231n: Deep Learning for Computer Vision.
<http://cs231n.stanford.edu/>

- [10] Ashwin Rao, Stanford CME241: Foundations of Reinforcement Learning with Applications in Finance.
<https://web.stanford.edu/class/cme241/>

- [11] Jimmy Ba & Bo Wang, UofT CSC413/2516: Neural Networks and Deep Learning.
<https://uoft-csc413.github.io/2022/>

Credits

Dr. Prashant Aparajeya

Computer Vision Scientist — Director(AISimply Ltd)

p.aparajeya@aisimply.uk

This project benefited from external collaboration, and we acknowledge their contribution with gratitude.