

Reinforcement Learning

Naeemullah Khan


naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning



Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, Sergey Levine

- Number of times the agent must interact with the environment in order to learn a task
- Good sample complexity is the first prerequisite for successful skill acquisition.
- Learning skills in the real world can take a substantial amount of time
 - can get damaged through trial and error

Main Problem: Sample Inefficiency


- "Learning Hand-Eye Coordination for Robotic Grasping with Deep Learning and Large-Scale Data Collection", Levine et al., 2016
 - 14 robot arms learning to grasp in parallel
 - objects started being picked up at around 20,000 grasps



<https://spectrum.ieee.org/automaton/robotics/artificial-intelligence/google-large-scale-robotic-grasping-project>

Main Problem: Sample Inefficiency

Observing the behavior of the robot after over 800,000 grasp attempts, which is equivalent to about 3000 robot-hours of practice, we can see the beginnings of intelligent reactive behaviors



https://www.youtube.com/watch?v=cXaic_k80uM

Main Problem: Sample Inefficiency

- Solution?
- Off-Policy Learning!

- On-policy learning: use the deterministic outcomes or samples from the target policy to train the algorithm
 - has **low sample efficiency** (TRPO, PPO, A3C)
 - require new samples to be collected for nearly every update to the policy
 - becomes extremely expensive when the task is complex
- Off-policy methods: training on a distribution of transitions or episodes produced by a different behavior policy rather than that produced by the target policy
 - Does not require full trajectories and can **reuse any past episodes** (experience replay) for much better sample efficiency
 - relatively straightforward for Q-learning based methods

- Value Function: How good is a state?

$$\begin{aligned} V(s) &= \mathbb{E}[G_t | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \dots) | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= \mathbb{E}[\underbrace{R_{t+1} + \gamma V(S_{t+1})}_{\text{temporal difference target}} | S_t = s] \end{aligned}$$

- Similarly, for Q-Function: How good is a state-action pair?

$$\begin{aligned} Q(s, a) &= \mathbb{E}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s, A_t = a] \\ &= \mathbb{E}[R_{t+1} + \gamma \mathbb{E}_{a \sim \pi} Q(S_{t+1}, a) | S_t = s, A_t = a] \end{aligned}$$

- $\dots, S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}, \dots$ (on-policy):

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$


- Q-Learning (off-policy)

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(S_{t+1}, a) - Q(S_t, A_t))$$

- DQN, Minh et al., 2015

$$\mathcal{L}(\theta) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right)^2 \right]$$

- Function Approximation
- Experience Replay: samples randomly drawn from replay memory
- Doesn't scale to continuous action space



- **Critic:** updates value function parameters w and depending on the algorithm it could be action-value $Q(a|s; w)$ or state-value $V(s; w)$.
- **Actor:** updates policy parameters θ , in the direction suggested by the critic, $\pi(a|s; \theta)$.


$$\nabla J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla \ln \pi(a|s, \theta) Q_{\pi}(s, a)] \quad \text{policy gradient}$$

$$\theta \leftarrow \theta + \alpha_{\theta} Q(s, a; w) \nabla_{\theta} \ln \pi(a|s; \theta) \quad \text{update actor}$$



$$G_{t:t+1} = r_t + \gamma Q(s', a'; w) - Q(s, a; w) \quad \text{correction for action-value}$$

$$w \leftarrow w + \alpha_w G_{t:t+1} \nabla_w Q(s, a; w). \quad \text{update critic}$$



- DDPG = DQN + DPG (Lillicrap et al., 2015)
 - off-policy actor-critic method that learns a deterministic policy in continuous domain
 - exploration noise added to the deterministic policy when select action
 - difficult to stabilize and brittle to hyperparameters (Duan et al., 2016, Henderson et al., 2017)
 - unscalable to complex tasks with high dimensions (Gu et al., 2017)




- Training is sensitive to randomness in the environment, initialization of the policy and the algorithm implementation




- Knowing only one way to act makes agents vulnerable to environmental changes that are common in the real-world





Traditional Graph of MDP



Graphical Model with Optimality Variables

Normal trajectory distribution

$$p(\tau) = p(\mathbf{s}_1, \mathbf{a}_1, \dots, \mathbf{s}_T, \mathbf{a}_T | \theta) = p(\mathbf{s}_1) \prod_{t=1}^T p(\mathbf{a}_t | \mathbf{s}_t, \theta) p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t).$$

Posterior trajectory distribution

$$p(\mathcal{O}_t = 1 | \mathbf{s}_t, \mathbf{a}_t) = \exp(r(\mathbf{s}_t, \mathbf{a}_t)).$$

$$\begin{aligned} p(\tau | \mathbf{o}_{1:T}) &\propto p(\tau, \mathbf{o}_{1:T}) = p(\mathbf{s}_1) \prod_{t=1}^T p(\mathcal{O}_t = 1 | \mathbf{s}_t, \mathbf{a}_t) p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t) \\ &= p(\mathbf{s}_1) \prod_{t=1}^T \exp(r(\mathbf{s}_t, \mathbf{a}_t)) p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t) \\ &= \left[p(\mathbf{s}_1) \prod_{t=1}^T p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t) \right] \exp \left(\sum_{t=1}^T r(\mathbf{s}_t, \mathbf{a}_t) \right). \end{aligned}$$

Variational Inference

$$\begin{aligned} D_{\text{KL}}(\hat{p}(\tau) \| p(\tau)) &= -E_{\tau \sim \hat{p}(\tau)} [\log p(\tau) - \log \hat{p}(\tau)]. \\ -D_{\text{KL}}(\hat{p}(\tau) \| p(\tau)) &= E_{\tau \sim \hat{p}(\tau)} \left[\log p(\mathbf{s}_1) + \sum_{t=1}^T (\log p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t) + r(\mathbf{s}_t, \mathbf{a}_t)) - \right. \\ &\quad \left. \log p(\mathbf{s}_1) - \sum_{t=1}^T (\log p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t) + \log \pi(\mathbf{a}_t | \mathbf{s}_t)) \right] \\ &= E_{\tau \sim \hat{p}(\tau)} \left[\sum_{t=1}^T r(\mathbf{s}_t, \mathbf{a}_t) - \log \pi(\mathbf{a}_t | \mathbf{s}_t) \right] \\ &= \sum_{t=1}^T E_{(\mathbf{s}_t, \mathbf{a}_t) \sim \hat{p}(\mathbf{s}_t, \mathbf{a}_t)} [r(\mathbf{s}_t, \mathbf{a}_t) - \log \pi(\mathbf{a}_t | \mathbf{s}_t)] \\ &= \sum_{t=1}^T E_{(\mathbf{s}_t, \mathbf{a}_t) \sim \hat{p}(\mathbf{s}_t, \mathbf{a}_t)} [r(\mathbf{s}_t, \mathbf{a}_t)] + E_{\mathbf{s}_t \sim \hat{p}(\mathbf{s}_t)} [\mathcal{H}(\pi(\mathbf{a}_t | \mathbf{s}_t))]. \end{aligned}$$

Conventional RL Objective - Expected Reward

$$\sum_t \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} [r(\mathbf{s}_t, \mathbf{a}_t)]$$





Maximum Entropy RL Objective - Expected Reward + Entropy of Policy



$$\sum_t \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} [r(\mathbf{s}_t, \mathbf{a}_t) + \alpha \mathcal{H}(\pi(\cdot | \mathbf{s}_t))]$$

Entropy of a RV x

$$H(P) = \mathbb{E}_{x \sim P} [-\log P(x)]$$

- MaxEnt RL agent can capture different modes of optimality to improve robustness against environmental changes





$$\min_{\pi} D_{KL}[\pi(\cdot | s_0) || \exp(Q(s_0, \cdot))]$$

$$= \min_{\pi} E_{\pi} \left[\log \frac{\pi(a_0 | s_0)}{\exp(Q(s_0, a_0))} \right]$$

$$= \max_{\pi} E_{\pi} [Q(s_0, a_0) - \log \pi(a_0 | s_0)]$$

$$= \max_{\pi} E_{\pi} \left[\sum_t r(s_t, a_t) + \mathcal{H}(\pi(\cdot | s_0)) | s_0 \right]$$

$$= J_{MaxEnt}(\pi(\cdot | s_0))$$

- Soft Q-Learning (Haarnoja et al., 2017)
 - off-policy algorithms under MaxEnt RL objective
 - Learns Q^* directly
 - Instead of taking max action, sample the action
 - use approximate inference methods to sample
 - Stein variational gradient descent
 - not true actor-critic

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(S_{t+1}, a) - Q(S_t, A_t))$$

$$Q_{\text{soft}}(\mathbf{s}_t, \mathbf{a}_t) \leftarrow r_t + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p_s} [V_{\text{soft}}(\mathbf{s}_{t+1})], \quad \forall \mathbf{s}_t, \mathbf{a}_t$$

$$V_{\text{soft}}(\mathbf{s}_t) \leftarrow \alpha \log \int_{\mathcal{A}} \exp \left(\frac{1}{\alpha} Q_{\text{soft}}(\mathbf{s}_t, \mathbf{a}') \right) d\mathbf{a}', \quad \forall \mathbf{s}_t$$

- One of the most efficient model-free algorithms
 - SOTA off-policy
 - well suited for real world robotics learning
- Can learn stochastic policy on continuous action domain
- Robust to noise
- Ingredients:
 - Actor-critic architecture with separate policy and value function networks
 - Off-policy formulation to reuse of previously collected data for efficiency
 - Entropy-constrained objective to encourage stability and exploration

- policy evaluation: compute value of π according to Max Entropy RL Objective
- modified Bellman backup operator \mathcal{T} :

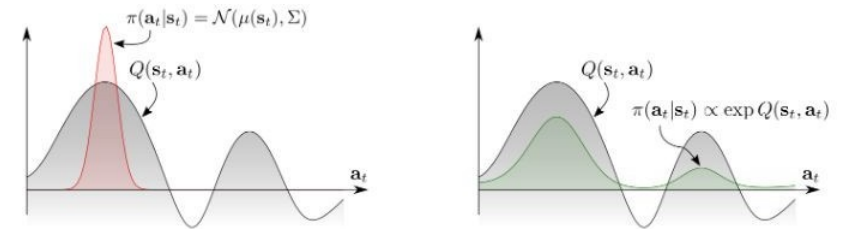
$$\mathcal{T}^\pi Q(\mathbf{s}_t, \mathbf{a}_t) \triangleq r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p} [V(\mathbf{s}_{t+1})]$$

$$V(\mathbf{s}_t) = \mathbb{E}_{\mathbf{a}_t \sim \pi} [Q(\mathbf{s}_t, \mathbf{a}_t) - \alpha \log \pi(\mathbf{a}_t | \mathbf{s}_t)]$$

- Lemma 1: Contraction Mapping for Soft Bellman Updates

$$Q^{k+1} = \mathcal{T}^\pi Q^k \quad \text{converges to the soft Q-function of } \pi$$

- policy improvement: update policy towards the exponential of the new soft Q-function
- modified Bellman backup operator T:
 - choose tractable family of distributions big Π
 - choose KL divergence to project the improved policy into big Π

$$\pi_{\text{new}} = \arg \min_{\pi' \in \Pi} D_{\text{KL}} \left(\pi'(\cdot | \mathbf{s}_t) \parallel \frac{\exp \left(\frac{1}{\alpha} Q^{\pi_{\text{old}}}(\mathbf{s}_t, \cdot) \right)}{Z^{\pi_{\text{old}}}(\mathbf{s}_t)} \right)$$


- Lemma 2

$$Q^{\pi_{\text{new}}}(\mathbf{s}_t, \mathbf{a}_t) \geq Q^{\pi_{\text{old}}}(\mathbf{s}_t, \mathbf{a}_t) \quad \text{for any state action pair}$$

- soft policy iteration: soft policy evaluation \leftrightarrow soft policy improvement
- Theorem 1: Repeated application of soft policy evaluation and soft policy improvement from any policy $\pi \in \Pi$ converges to the optimal MaxEnt policy among all policies in Π
 - exact form applicable only in discrete case
 - need function approximation to represent Q-values in continuous domains
 - \rightarrow Soft Actor-Critic (SAC)!

$$Q_{\theta}(\mathbf{s}_t, \mathbf{a}_t)$$

- e.g. neural network

$$\pi_{\phi}(\mathbf{a}_t | \mathbf{s}_t)$$

parameterized tractable policy

- e.g. Gaussian with mean and covariances given by neural networks

$$J_Q(\theta) \quad \hat{\nabla}_{\theta} J_Q(\theta)$$

soft Q-function objective and its stochastic gradient wrt its parameters

$$J_{\pi}(\phi) \quad \hat{\nabla}_{\phi} J_{\pi}(\phi)$$

policy objective and stochastic gradient wrt its parameters

- Critic - Soft Q-function
 - minimize square error
 - $\bar{\theta}$ exponential moving average of soft Q-function weights to stabilize training (DQN)

$$J_Q(\theta) = \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \mathcal{D}} \left[\frac{1}{2} \left(Q_{\theta}(\mathbf{s}_t, \mathbf{a}_t) - \left(r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p} [V_{\bar{\theta}}(\mathbf{s}_{t+1})] \right) \right)^2 \right]$$

$$V(\mathbf{s}_t) = \mathbb{E}_{\mathbf{a}_t \sim \pi} [Q(\mathbf{s}_t, \mathbf{a}_t) - \alpha \log \pi(\mathbf{a}_t | \mathbf{s}_t)]$$

$$\hat{\nabla}_{\theta} J_Q(\theta) = \nabla_{\theta} Q_{\theta}(\mathbf{a}_t, \mathbf{s}_t) (Q_{\theta}(\mathbf{s}_t, \mathbf{a}_t) - (r(\mathbf{s}_t, \mathbf{a}_t) + \gamma (Q_{\bar{\theta}}(\mathbf{s}_{t+1}, \mathbf{a}_{t+1}) - \alpha \log (\pi_{\phi}(\mathbf{a}_{t+1} | \mathbf{s}_{t+1}))))$$

- Actor - Policy $\pi_{\text{new}} = \arg \min_{\pi' \in \Pi} D_{\text{KL}} \left(\pi'(\cdot | \mathbf{s}_t) \parallel \frac{\exp \left(\frac{1}{\alpha} Q^{\pi_{\text{old}}}(\mathbf{s}_t, \cdot) \right)}{Z^{\pi_{\text{old}}}(\mathbf{s}_t)} \right)$

- multiply by alpha and ignoring the normalization Z

$$J_{\pi}(\phi) = \mathbb{E}_{\mathbf{s}_t \sim \mathcal{D}} \left[\mathbb{E}_{\mathbf{a}_t \sim \pi_{\phi}} [\alpha \log (\pi_{\phi}(\mathbf{a}_t | \mathbf{s}_t)) - Q_{\theta}(\mathbf{s}_t, \mathbf{s}_t)] \right]$$

- reparameterize with neural network f $\mathbf{a}_t = f_{\phi}(\epsilon_t; \mathbf{s}_t)$
 - epsilon: input noise vector, sampled from a fixed distribution (spherical Gaussian)

$$J_{\pi}(\phi) = \mathbb{E}_{\mathbf{s}_t \sim \mathcal{D}, \epsilon_t \sim \mathcal{N}} [\alpha \log \pi_{\phi}(f_{\phi}(\epsilon_t; \mathbf{s}_t) | \mathbf{s}_t) - Q_{\theta}(\mathbf{s}_t, f_{\phi}(\epsilon_t; \mathbf{s}_t))]$$

$$\hat{\nabla}_{\phi} J_{\pi}(\phi) = \nabla_{\phi} \alpha \log (\pi_{\phi}(\mathbf{a}_t | \mathbf{s}_t)) + (\nabla_{\mathbf{a}_t} \alpha \log (\pi_{\phi}(\mathbf{a}_t | \mathbf{s}_t)) - \nabla_{\mathbf{a}_t} Q(\mathbf{s}_t, \mathbf{a}_t)) \nabla_{\phi} f_{\phi}(\epsilon_t; \mathbf{s}_t)$$

- Unbiased gradient estimator that extends DDPG stype policy gradients to any tractable stochastic policy

Algorithm 1 Soft Actor-Critic

- 1: Input: initial policy parameters θ , Q-function parameters ϕ_1, ϕ_2 , empty replay buffer \mathcal{D}
- 2: Set target parameters equal to main parameters $\phi_{\text{targ},1} \leftarrow \phi_1, \phi_{\text{targ},2} \leftarrow \phi_2$
- 3: **repeat**
- 4: Observe state s and select action $a \sim \pi_\theta(\cdot|s)$
- 5: Execute a in the environment
- 6: Observe next state s' , reward r , and done signal d to indicate whether s' is terminal
- 7: Store (s, a, r, s', d) in replay buffer \mathcal{D}
- 8: If s' is terminal, reset environment state.
- 9: **if** it's time to update **then**
- 10: **for** j in range(however many updates) **do**
- 11: Randomly sample a batch of transitions, $B = \{(s, a, r, s', d)\}$ from \mathcal{D}
- 12: Compute targets for the Q functions:

$$y(r, s', d) = r + \gamma(1 - d) \left(\min_{i=1,2} Q_{\phi_{\text{targ},i}}(s', \tilde{a}') - \alpha \log \pi_\theta(\tilde{a}'|s') \right), \quad \tilde{a}' \sim \pi_\theta(\cdot|s')$$

- 13: Update Q-functions by one step of gradient descent using

$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi_i}(s, a) - y(r, s', d))^2 \quad \text{for } i = 1, 2$$

- 14: Update policy by one step of gradient ascent using

$$\nabla_\theta \frac{1}{|B|} \sum_{s \in B} \left(\min_{i=1,2} Q_{\phi_i}(s, \tilde{a}_\theta(s)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s)|s) \right),$$


where $\tilde{a}_\theta(s)$ is a sample from $\pi_\theta(\cdot|s)$ which is differentiable wrt θ via the reparametrization trick.

- 15: Update target networks with


$$\phi_{\text{targ},i} \leftarrow \rho \phi_{\text{targ},i} + (1 - \rho) \phi_i \quad \text{for } i = 1, 2$$

- 16: **end for**
- 17: **end if**
- 18: **until** convergence

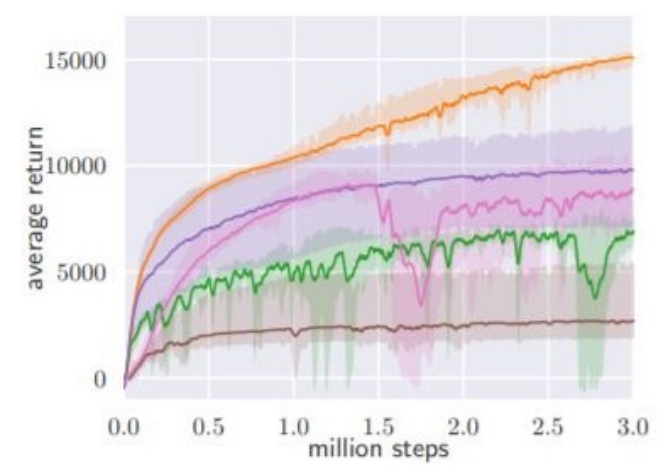
- Tasks
 - A range of continuous control tasks from the OpenAI gym benchmark suite
 - RL-Lab implementation of the Humanoid task
 - The easier tasks can be solved by a wide range of different algorithms, the more complex benchmarks, such as the 21-dimensional Humanoid (rllab) are exceptionally difficult to solve with off-policy algorithms.
- Baselines:
 - DDPG, SQL, PPO, TD3 (concurrent)
 - TD3 is an extension to DDPG that first applied the double Q-learning trick to continuous control along with other improvements.




(a) Hopper-v1



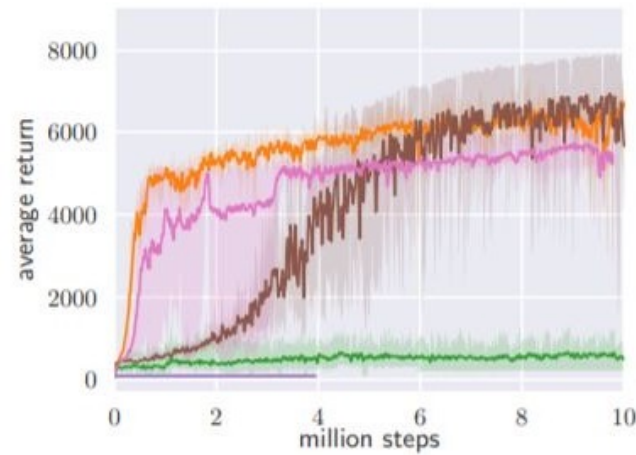
(b) Walker2d-v1




(c) HalfCheetah-v1



(d) Ant-v1

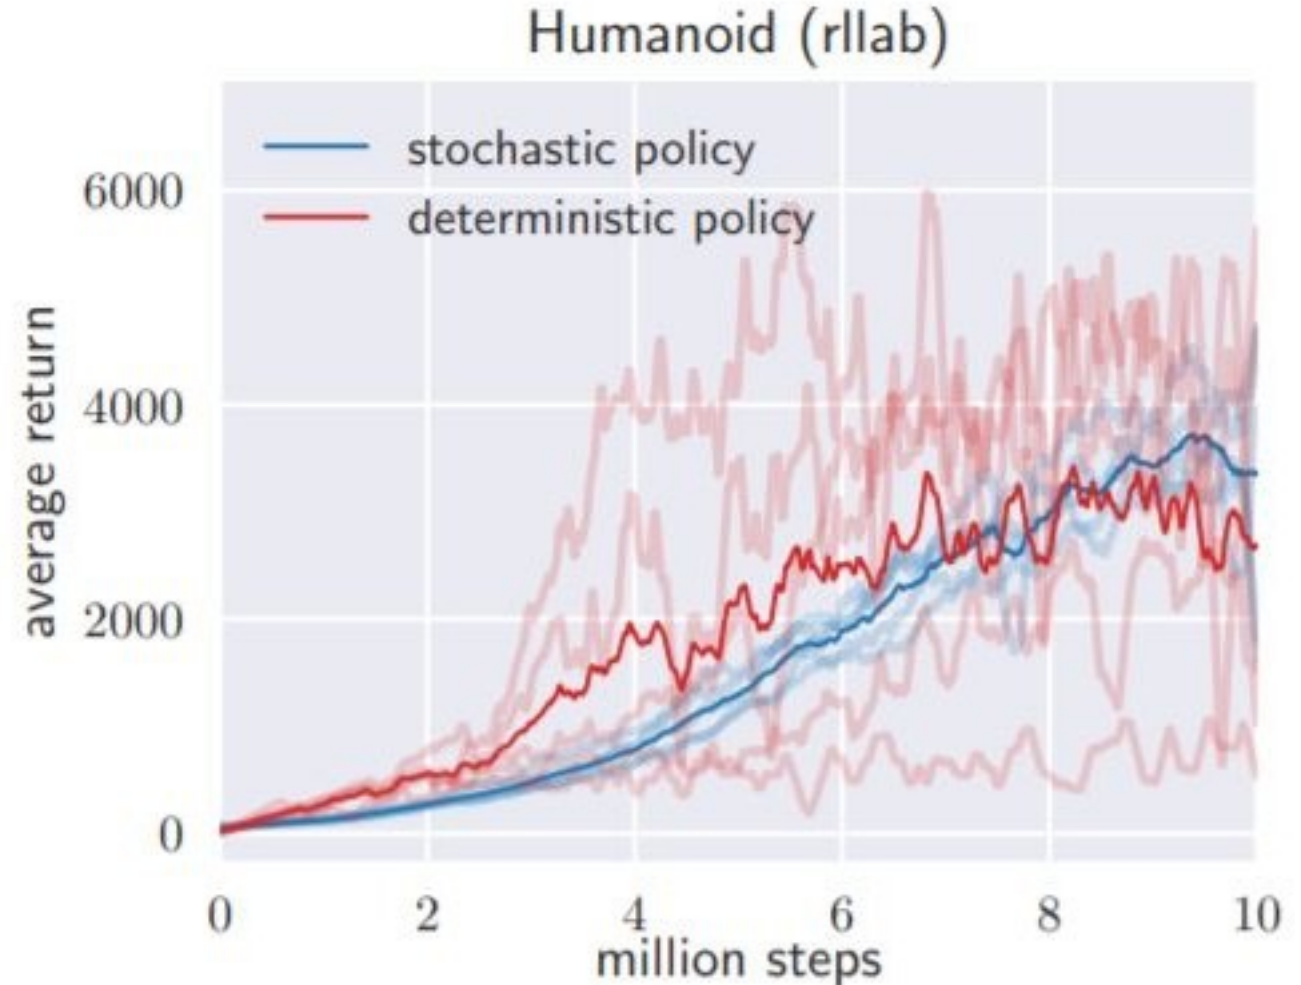


(e) Humanoid-v1




(f) Humanoid (rllab)

- How does the stochasticity of the policy and entropy maximization affect the performance?
- Comparison with a deterministic variant of SAC that does not maximize the entropy and that closely resembles DDPG




<https://arxiv.org/abs/1801.01290>


Experimental Results: Hyperparameter Sensitivity



(a) Evaluation



(b) Reward Scale



(c) Target Smoothing Coefficient (τ)


<https://arxiv.org/abs/1801.01290>

- Unfortunately, SAC also suffers from brittleness to the alpha temperature hyperparameter that controls exploration
 - -> automatic temperature tuning!


- Adaptive temperature coefficient

$$\sum_t \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} [r(\mathbf{s}_t, \mathbf{a}_t) + \overset{\text{temperature}}{\alpha} \mathcal{H}(\pi(\cdot | \mathbf{s}_t))]$$

- Extend to real-world tasks such as locomotion for a quadrupedal robot and robotic manipulation with a dexterous hand



- Dexterous Hand Manipulations
- 20 hour end-to-end learning
- valve position as input: SAC 3 hours vs. PPO 7.4 hours



- Choosing the optimal temperature is non-trivial (tuned for each task)
- Constrained optimization problem:

$$\max \sum_t \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} [r(\mathbf{s}_t, \mathbf{a}_t) + \alpha \mathcal{H}(\pi(\cdot | \mathbf{s}_t))]$$

$$\max_{\pi_{0:T}} \mathbb{E}_{\rho_\pi} \left[\sum_{t=0}^T r(\mathbf{s}_t, \mathbf{a}_t) \right] \text{ s.t. } \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} [-\log(\pi_t(\mathbf{a}_t | \mathbf{s}_t))] \geq \mathcal{H} \quad \forall t$$

Unroll the expectation

$$\max_{\pi_0} \left(\mathbb{E} [r(\mathbf{s}_0, \mathbf{a}_0)] + \max_{\pi_1} \left(\mathbb{E} [\dots] + \max_{\pi_T} \mathbb{E} [r(\mathbf{s}_T, \mathbf{a}_T)] \right) \right)$$

For the last time step in the trajectory

$$\max_{\pi_T} \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} [r(\mathbf{s}_T, \mathbf{a}_T)] = \min_{\alpha_T \geq 0} \max_{\pi_T} \mathbb{E} [r(\mathbf{s}_T, \mathbf{a}_T) - \alpha_T \log \pi(\mathbf{a}_T | \mathbf{s}_T)] - \alpha_T \mathcal{H},$$

$$\arg \min_{\alpha_T} \mathbb{E}_{\mathbf{s}_t, \mathbf{a}_t \sim \pi_t^*} [-\alpha_T \log \pi_T^*(\mathbf{a}_T | \mathbf{s}_T; \alpha_T) - \alpha_T \mathcal{H}].$$

Similarly, for the previous time step

$$\begin{aligned} & \max_{\pi_{T-1}} \left(\mathbb{E} [r(\mathbf{s}_{T-1}, \mathbf{a}_{T-1})] + \max_{\pi_T} \mathbb{E} [r(\mathbf{s}_T, \mathbf{a}_T)] \right) \\ &= \max_{\pi_{T-1}} \left(Q_{T-1}^*(\mathbf{s}_{T-1}, \mathbf{a}_{T-1}) - \alpha_T \mathcal{H} \right) \\ &= \min_{\alpha_{T-1} \geq 0} \max_{\pi_{T-1}} \left(\mathbb{E} [Q_{T-1}^*(\mathbf{s}_{T-1}, \mathbf{a}_{T-1})] - \mathbb{E} [\alpha_{T-1} \log \pi(\mathbf{a}_{T-1} | \mathbf{s}_{T-1})] - \alpha_{T-1} \mathcal{H} \right) + \alpha_T^* \mathcal{H}. \end{aligned} \tag{16}$$

$$\alpha_t^* = \arg \min_{\alpha_t} \mathbb{E}_{\mathbf{a}_t \sim \pi_t^*} \left[-\alpha_t \log \pi_t^*(\mathbf{a}_t | \mathbf{s}_t; \alpha_t) - \alpha_t \bar{\mathcal{H}} \right]$$

Algorithm 1 Soft Actor-Critic

Input: θ_1, θ_2, ϕ two soft Q-functions

$\bar{\theta}_1 \leftarrow \theta_1, \bar{\theta}_2 \leftarrow \theta_2$

$\mathcal{D} \leftarrow \emptyset$

for each iteration **do**

for each environment step **do**

$\mathbf{a}_t \sim \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)$

$\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$

$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$

end for

for each gradient step **do**

$\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i)$ for $i \in \{1, 2\}$

$\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi)$

$\alpha \leftarrow \alpha - \lambda \hat{\nabla}_\alpha J(\alpha)$

$\bar{\theta}_i \leftarrow \tau \theta_i + (1 - \tau) \bar{\theta}_i$ for $i \in \{1, 2\}$

end for exponential moving average

end for

Output: θ_1, θ_2, ϕ

▷ Initial parameters

▷ Initialize target network weights

▷ Initialize an empty replay pool

▷ Sample action from the policy

▷ Sample transition from the environment

▷ Store the transition in the replay pool

▷ Update the Q-function parameters

▷ Update policy weights

▷ Adjust temperature

▷ Update target network weights

▷ Optimized parameters

Experimental Results: RL Lab

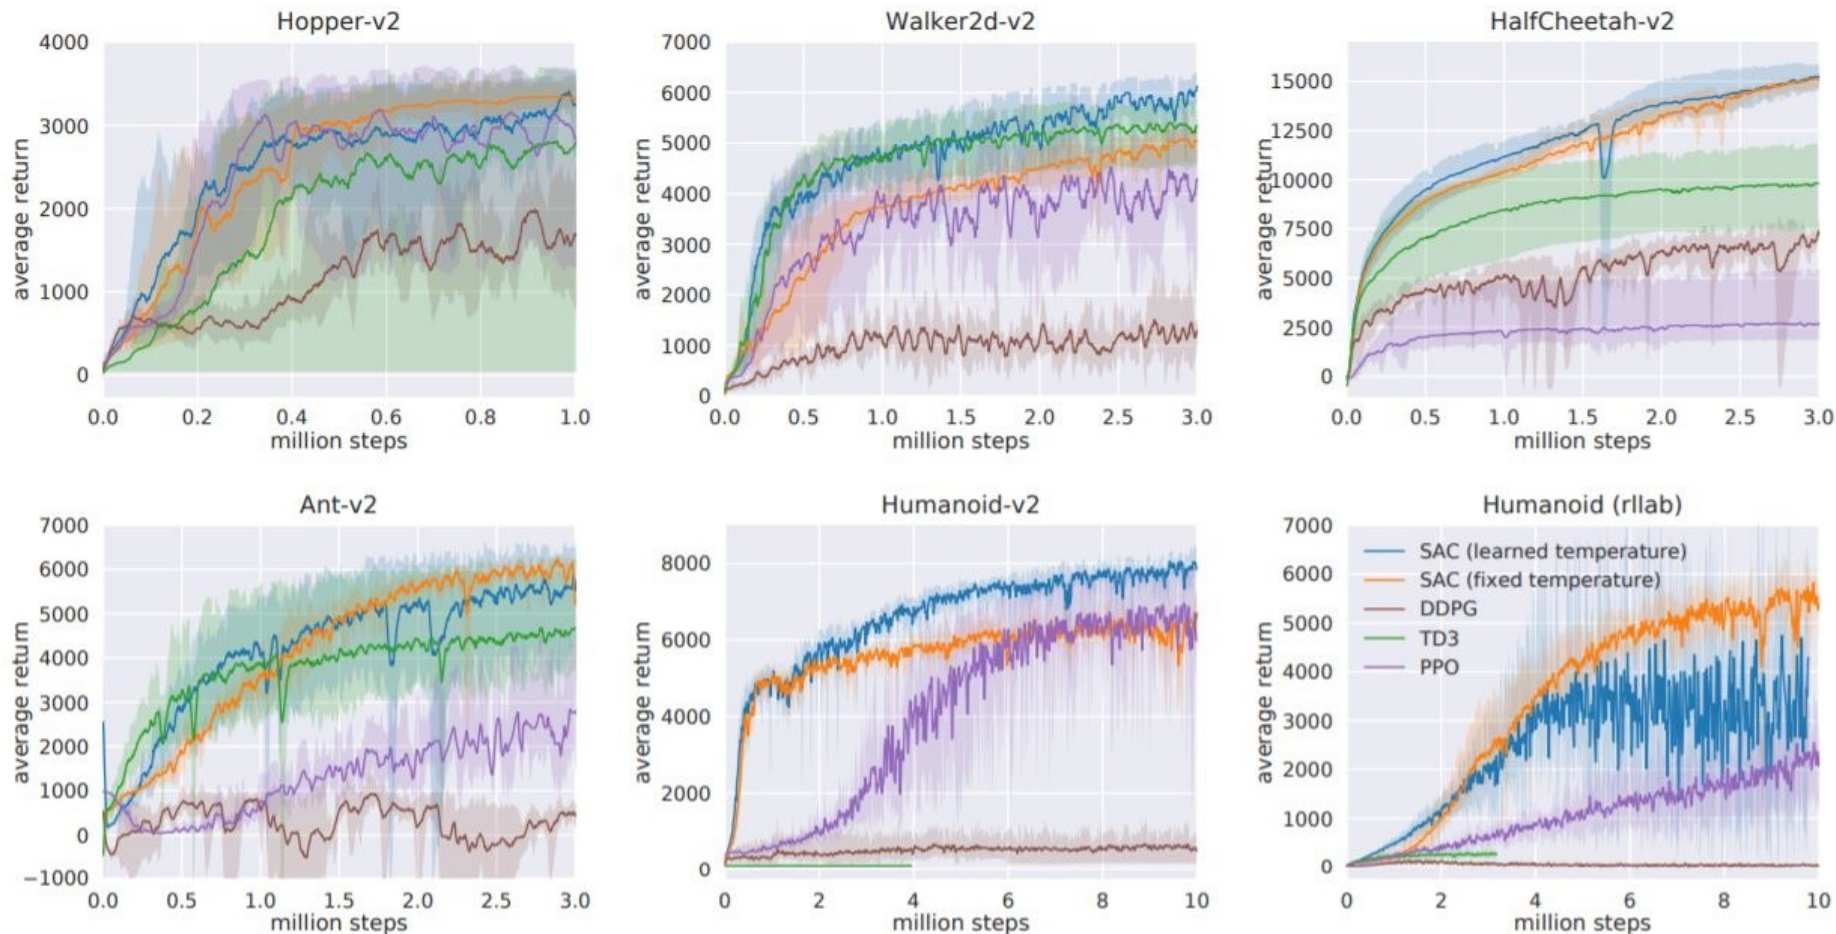




Figure 1: Training curves on continuous control benchmarks. Soft actor-critic (blue and yellow) performs consistently across all tasks and outperforming both on-policy and off-policy methods in the most challenging tasks.

Experimental Results: Robustness




- Lack of experiments on hard-exploration problem

- Lack of experiments on hard-exploration problem
- Approximating a multi-modal Boltzmann distribution with a unimodal Gaussian



- Lack of experiments on hard-exploration problem
- Approximating a multi-modal Boltzmann distribution with a unimodal Gaussian
- High-variance using automatic temperature tuning

- Lack of experiments on hard-exploration problem
- Approximating a multi-modal Boltzmann distribution with a unimodal Gaussian
- High-variance using automatic temperature tuning



- An off-policy maximum entropy deep reinforcement learning algorithm
 - Sample-efficient
 - Scale to high-dimensional observation/action space
 - Robustness to random seed, noise and etc.
- Theoretical Results
 - Convergence of soft policy iteration
 - Derivation of soft-actor critic algorithm
- Empirical Results
 - SAC outperforms SOTA model-free deep RL methods, including DDPG, PPO and Soft Q-learning, in terms of the policy's optimality, sample complexity and robustness.

These slides have been adapted from

- Animesh Garg, [CSC2621: Reinforcement Learning in Robotics, University of Toronto](#)