

# Large Language Models (LLMs)

Naeemullah Khan

[naeemullah.khan@kaust.edu.sa](mailto:naeemullah.khan@kaust.edu.sa)



جامعة الملك عبد الله  
للعلوم والتقنية  
King Abdullah University of  
Science and Technology



LMH  
Lady Margaret Hall

July 3, 2025

# Large Language Models

BCV

## Application Layer

### Copywriting



### Coding



### Dev Tools



### Chat / Comms



### BizOps



## Infrastructure Layer


### Model Creation



### Hardware



### Fine Tuning



### Inference



# Table of Contents

1. Motivation
2. Learning Outcomes
3. BERT and GPT Concepts
4. Scaling Laws for LLMs
5. Pre-training Overview
6. Tokenization: Why It Matters?
7. Level of Tokenization
8. Byte-Pair Encoding (BPE)
9. Token-Free LLMs
10. Context Window Problems
11. Moving Towards Larger Context Windows
12. Limitations of LLMs
13. Future Directions

## ► Why LLMs?

- Transforming NLP: ChatGPT, Claude, Gemini, etc.
- Achieving human-like generation and comprehension.
- Pivotal for tasks like summarization, translation, reasoning.

## ► Need for deeper understanding:

- LLMs are expensive to train and operate.
- Design decisions impact performance significantly (e.g., tokenization, scaling laws).



By the end of this session, you should be able to:

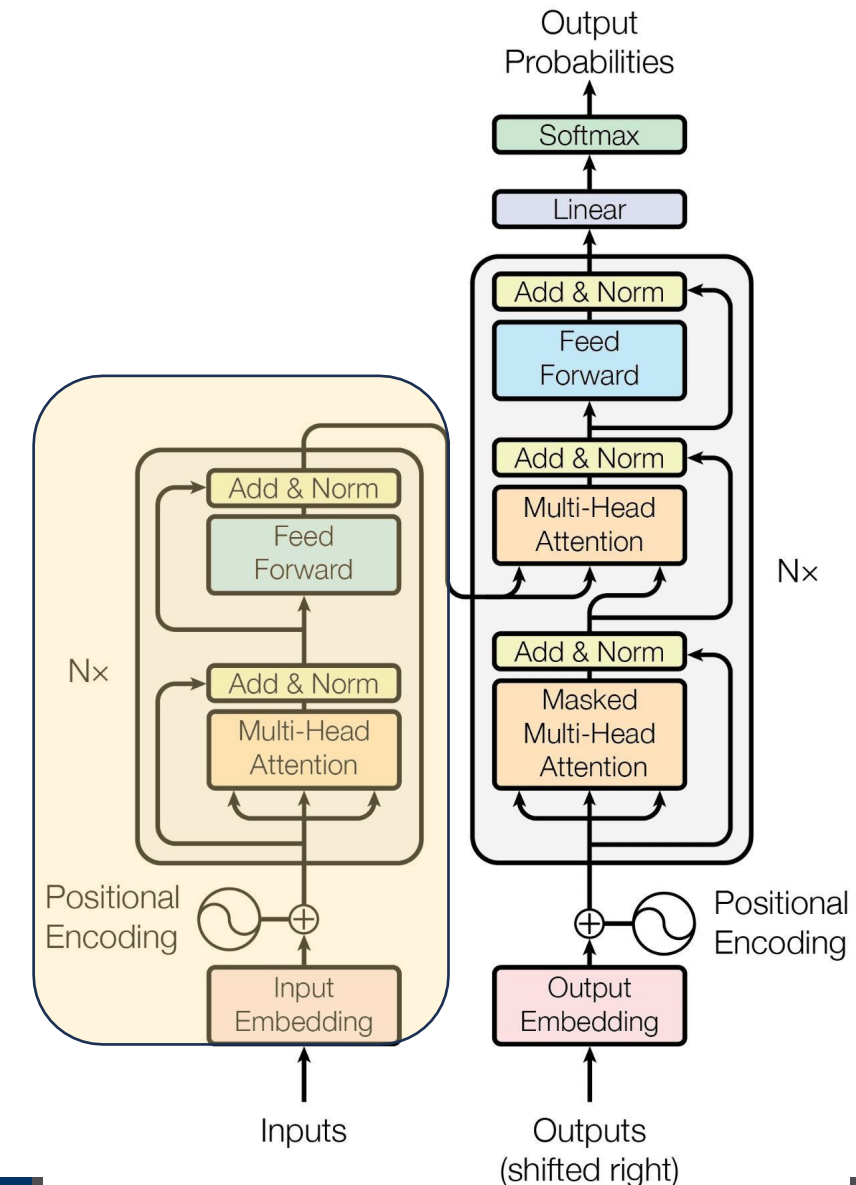
- ▶ Explain core architectures like BERT and GPT.
- ▶ Understand scaling laws for LLM development.
- ▶ Describe tokenization strategies and their impact.
- ▶ Discuss limitations of current models (e.g., context window).
- ▶ Explore emerging directions like token-free LLMs.

## ► BERT (Bidirectional Encoder Representations from Transformers)

- Uses transformer encoder only.
- Trained with Masked Language Modeling (MLM).
- **Bidirectional:** Considers both left and right context.
- Fine-tuned for tasks like QA, classification.
- **Architecture:**
  - Layers of encoder blocks.
  - Positional encodings added to embeddings.
  - Self-attention heads capture dependencies.



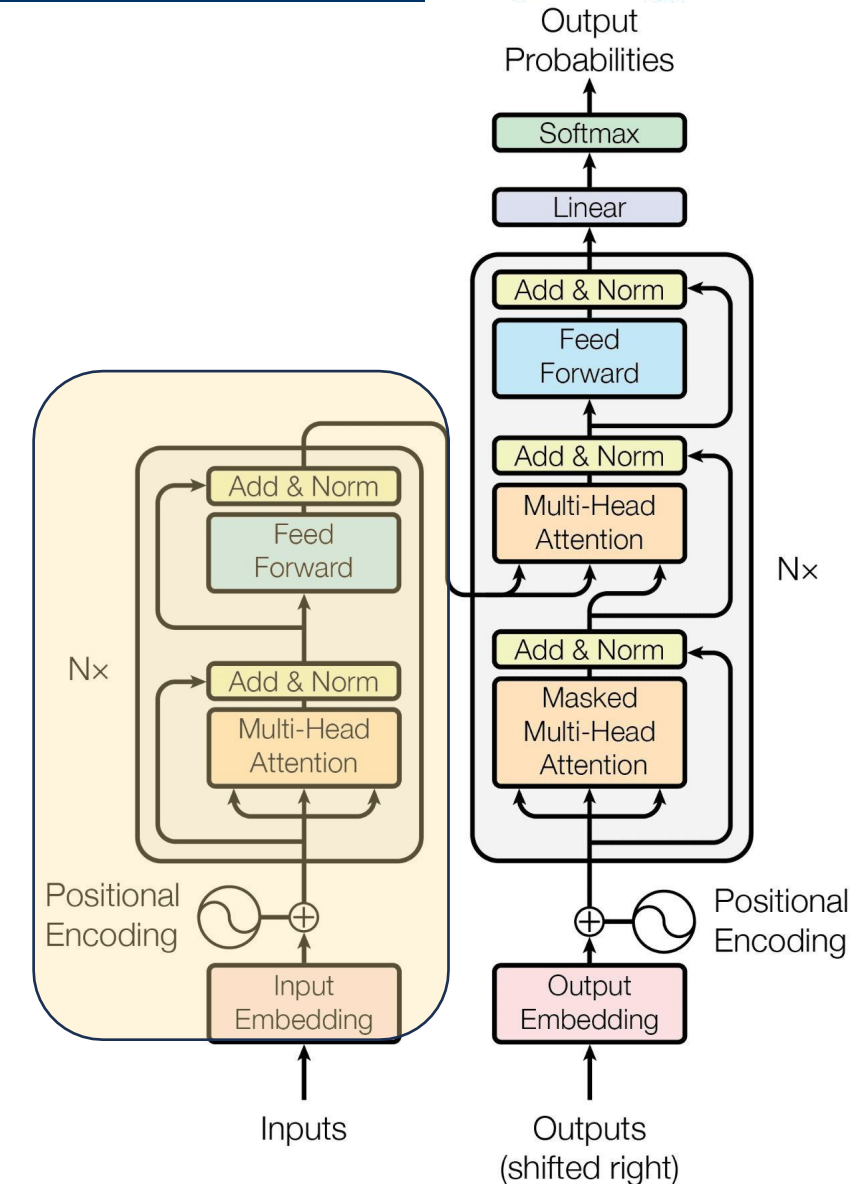
- One of the biggest challenges in LM-building used to be the lack of task-specific training data.
- What if we learn an effective representation that can be applied to a variety of downstream tasks?
  - Word2vec (2013)
  - GloVe (2014)





## BERT Pre-Training Corpus:

- English Wikipedia - 2,500 million words
- Book Corpus - 800 million words








## BERT Pre-Training Corpus:

- English Wikipedia - 2,500 million words
- Book Corpus - 800 million words

## BERT Pre-Training Tasks:

- MLM (Masked Language Modeling)
- NSP (Next Sentence Prediction)





## BERT Pre-Training Corpus:


- English Wikipedia - 2,500 million words
- Book Corpus - 800 million words

## BERT Pre-Training Tasks:

- MLM (Masked Language Modeling)
- NSP (Next Sentence Prediction)

## BERT Pre-Training Results:

- BERT-Base – 110M Params
- BERT-Large – 340M Params



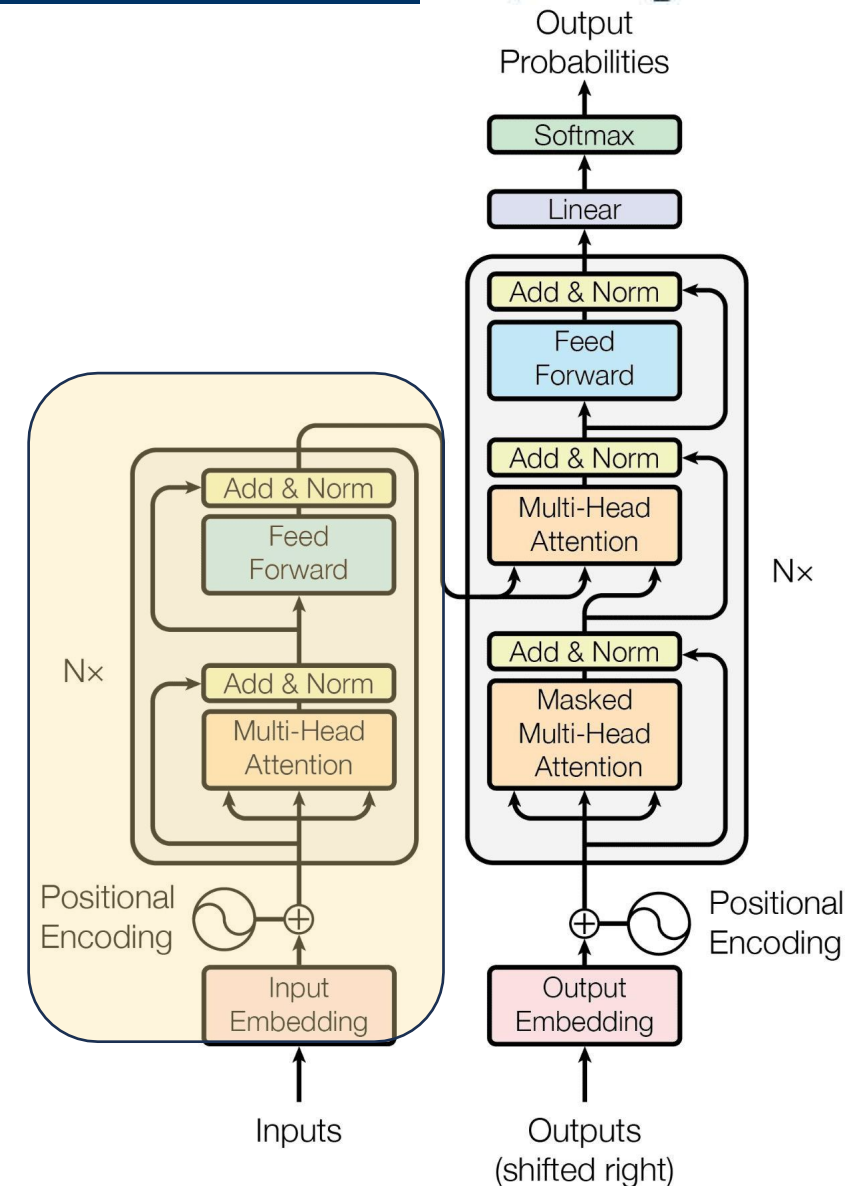
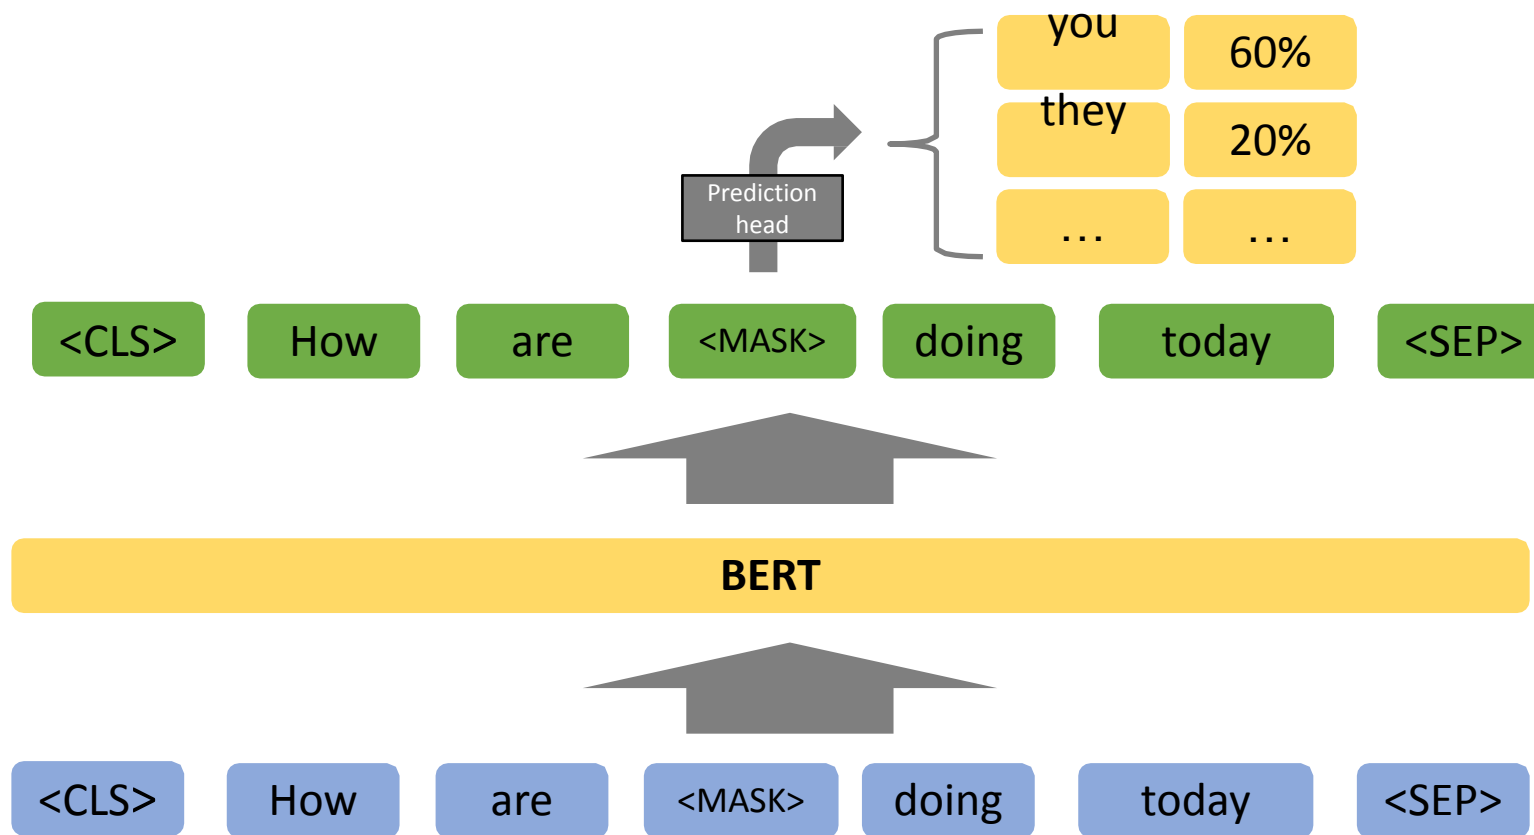
# BERT - Bidirectional Encoder Representations

أكاديمية كاوست  
KAUST ACADEMY



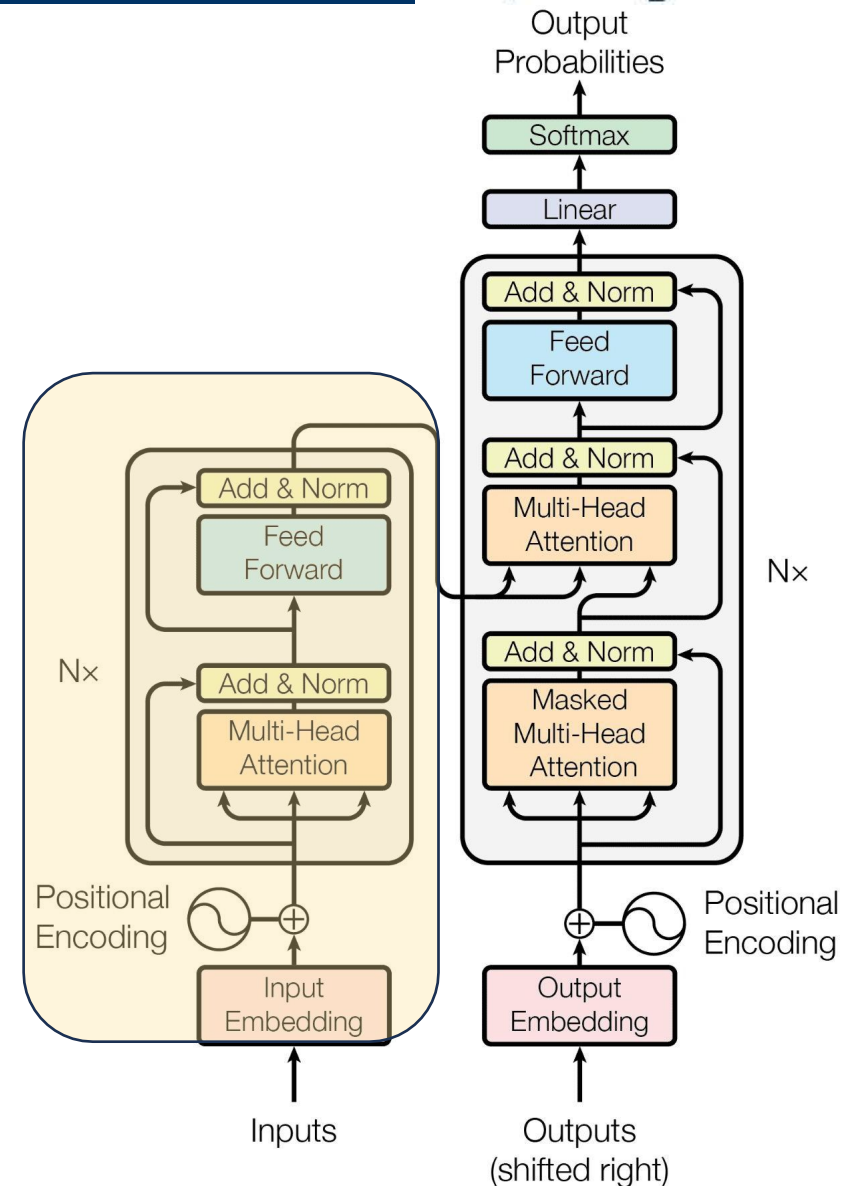
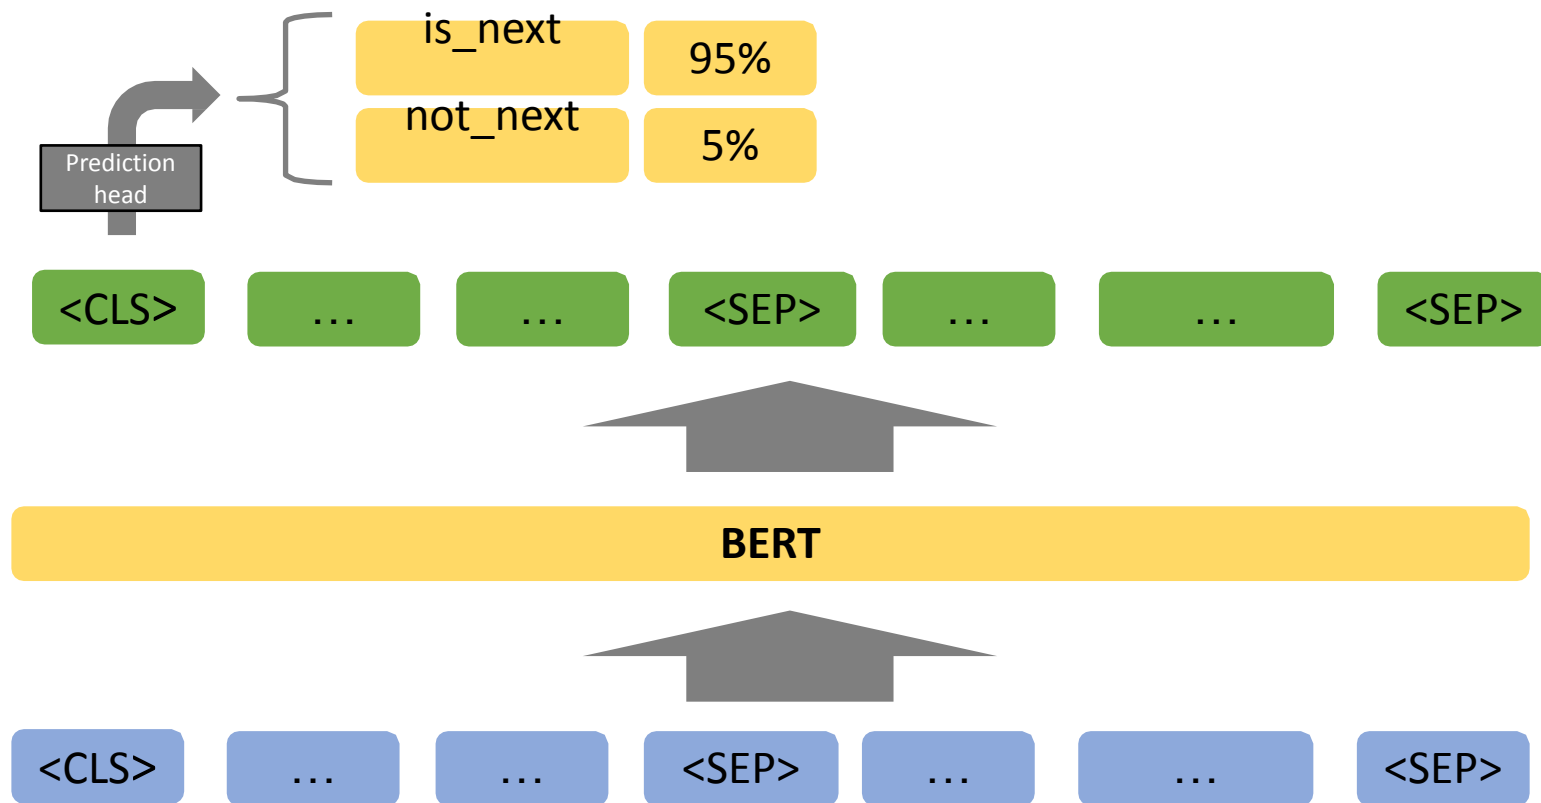
Lady Margaret Hall

## MLM (Masked Language Modeling)





## NSP (Next Sentence Prediction)





## BERT Fine-Tuning:

- Simply add a task-specific module after the last encoder layer to map it to the desired dimension.

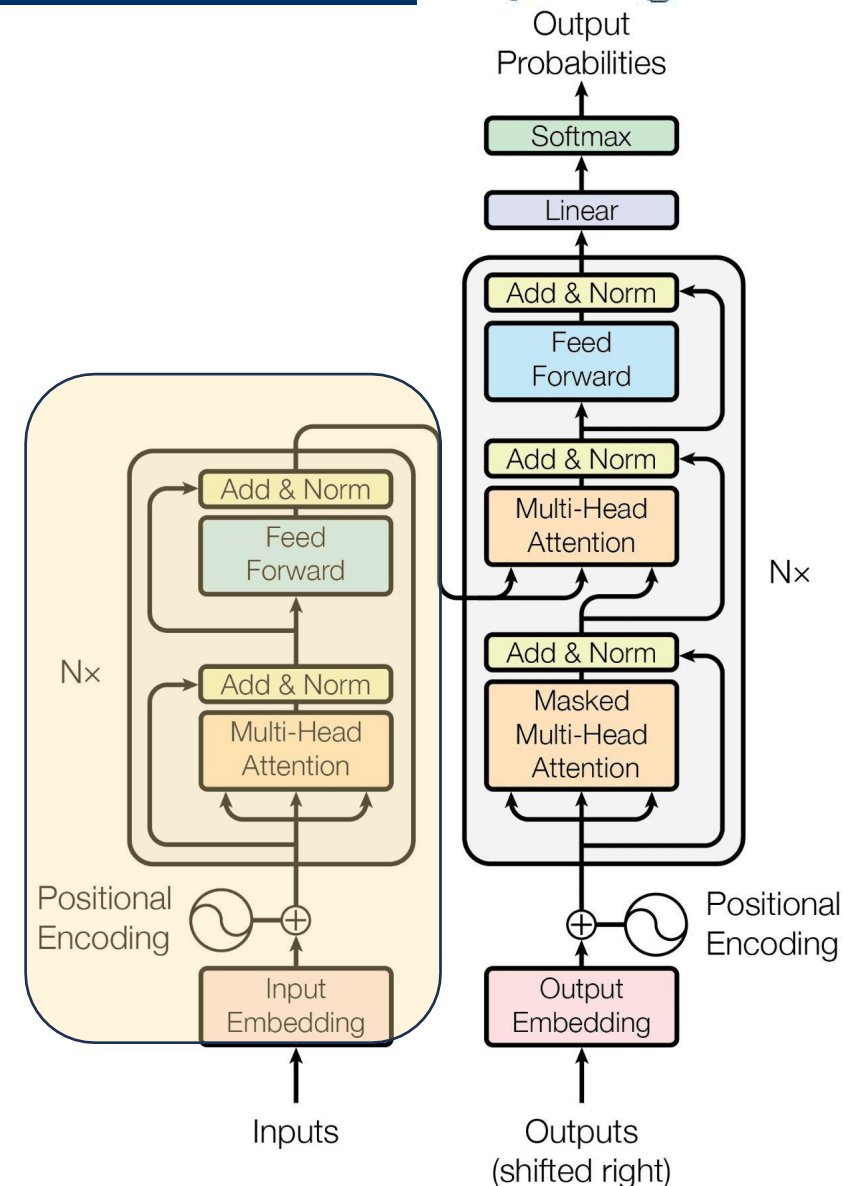
- Classification Tasks:

- Add a feed-forward layer on top of the encoder output for the [CLS] token

- Question Answering Tasks:

- Train two extra vectors to mark the beginning and end of answer from paragraph

- ...








## BERT Evaluation:

- General Language Understanding Evaluation (GLUE)
  - Sentence pair tasks
  - Single sentence classification
- Stanford Question Answering Dataset (SQuAD)



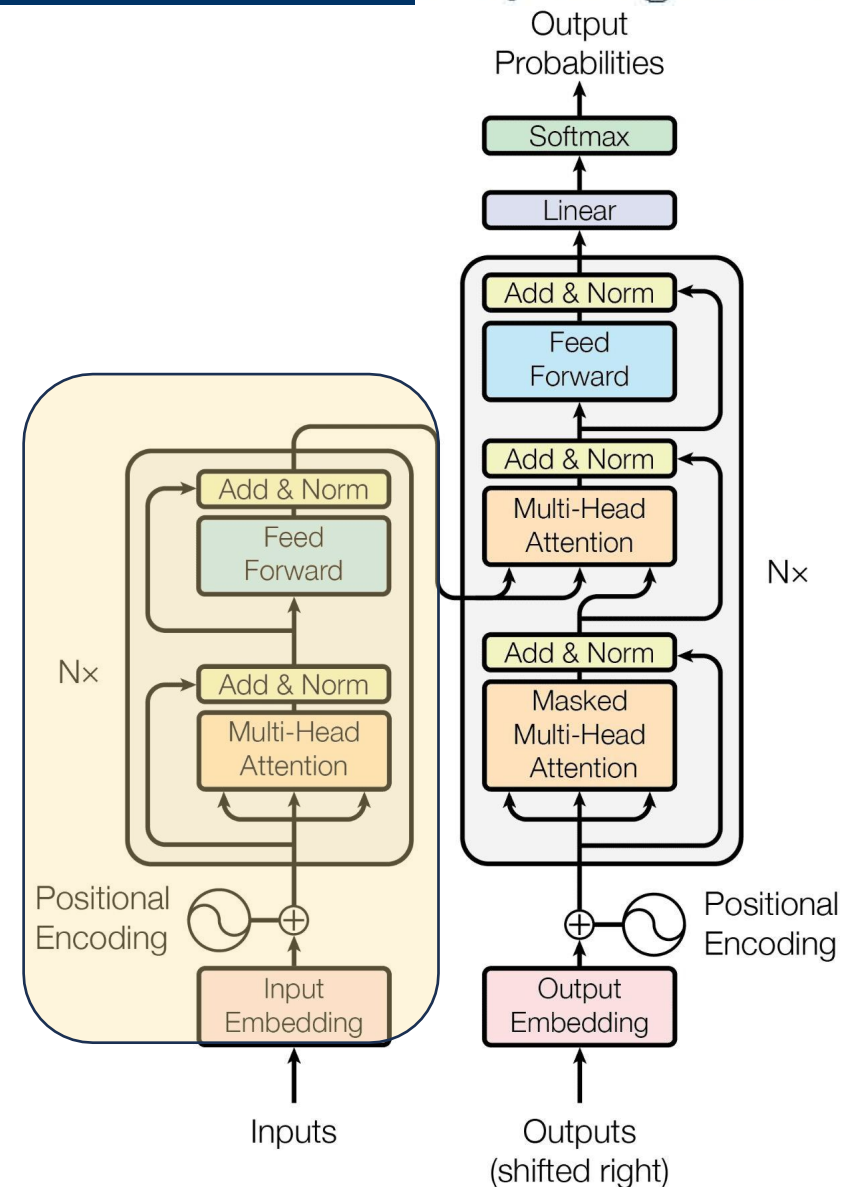


## BERT Evaluation:

System	MNLI-(m/mm) 392k	QQP 363k	QNLI 108k	SST-2 67k	CoLA 8.5k	STS-B 5.7k	MRPC 3.5k	RTE 2.5k	Average
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT <sub>BASE</sub>	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT <sub>LARGE</sub>	<b>86.7/85.9</b>	<b>72.1</b>	<b>92.7</b>	<b>94.9</b>	<b>60.5</b>	<b>86.5</b>	<b>89.3</b>	<b>70.1</b>	<b>82.1</b>

System	Dev EM	F1	Test EM	F1
Leaderboard (Oct 8th, 2018)				
Human	-	-	82.3	91.2
#1 Ensemble - nlnet	-	-	86.0	91.7
#2 Ensemble - QANet	-	-	84.5	90.5
#1 Single - nlnet	-	-	83.5	90.1
#2 Single - QANet	-	-	82.5	89.3
Published				
BiDAF+ELMo (Single)	-	85.8	-	-
R.M. Reader (Single)	78.9	86.3	79.5	86.6
R.M. Reader (Ensemble)	81.2	87.9	82.3	88.5
Ours				
BERT <sub>BASE</sub> (Single)	80.8	88.5	-	-
BERT <sub>LARGE</sub> (Single)	84.1	90.9	-	-
BERT <sub>LARGE</sub> (Ensemble)	85.8	91.8	-	-
BERT <sub>LARGE</sub> (Sgl.+TriviaQA)	<b>84.2</b>	<b>91.1</b>	<b>85.1</b>	<b>91.8</b>
BERT <sub>LARGE</sub> (Ens.+TriviaQA)	<b>86.2</b>	<b>92.2</b>	<b>87.4</b>	<b>93.2</b>


Table 2: SQuAD results. The BERT ensemble is 7x systems which use different pre-training checkpoints and fine-tuning seeds.





## What is our takeaway from BERT?


- **Pre-training tasks can be invented flexibly...**
  - Effective representations can be derived from a flexible regime of pre-training tasks.





## What is our takeaway from BERT?

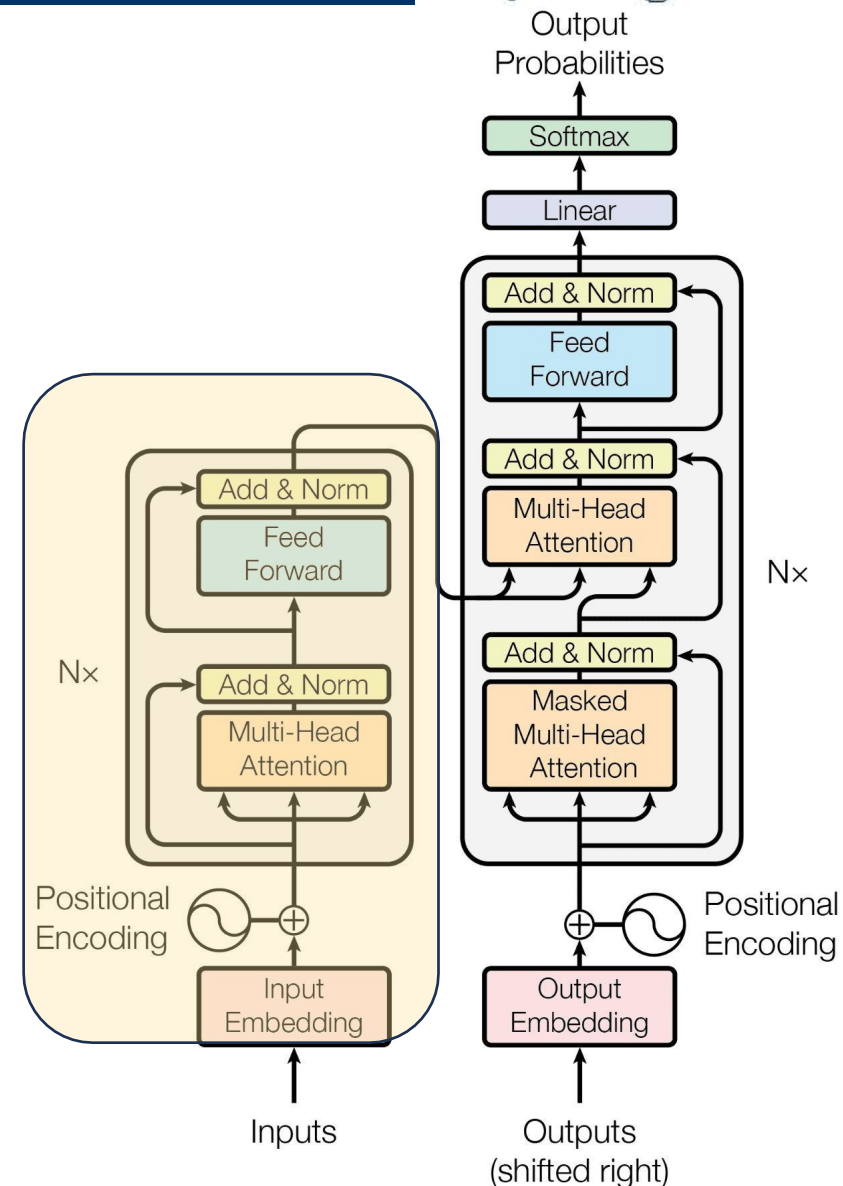
- **Pre-training tasks can be invented flexibly...**
  - Effective representations can be derived from a flexible regime of pre-training tasks.
- **Different NLP tasks seem to be highly transferable with each other...**
  - As long as we have effective representations, that seems to form a general model which can serve as the backbone for many specialized models.





## What is our takeaway from BERT?

- **Pre-training tasks can be invented flexibly...**
  - Effective representations can be derived from a flexible regime of pre-training tasks.
- **Different NLP tasks seem to be highly transferable with each other...**
  - As long as we have effective representations, that seems to form a general model which can serve as the backbone for many specialized models.
- **And scaling works!!!**
  - 340M was considered large in 2018



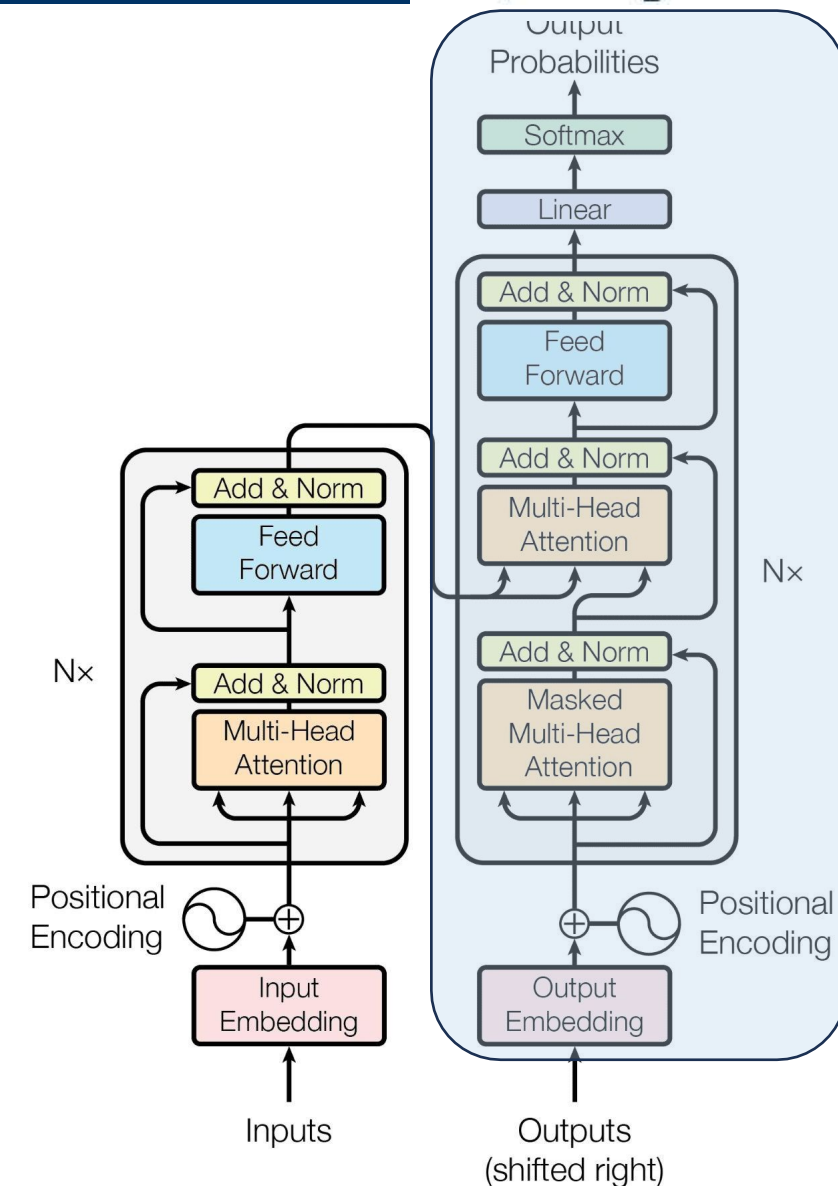


## ► GPT (Generative Pre-trained Transformer)

- Uses transformer decoder only.
- Trained with next-token prediction.
- **Unidirectional:** Considers left context only.
- Fine-tuned for text generation, dialogue systems.
- **Architecture:**
  - Layers of decoder blocks.
  - Causal masking to prevent future token access.
  - Self-attention heads capture sequential dependencies.

# GPT – **Generative** Pretrained Transformer

- Similarly motivated as BERT, though differently designed
- Can we leverage large amounts of unlabeled data to pretrain an LM that understands general patterns?



# GPT – **Generative** Pretrained Transformer

## GPT Pre-Training Corpus:

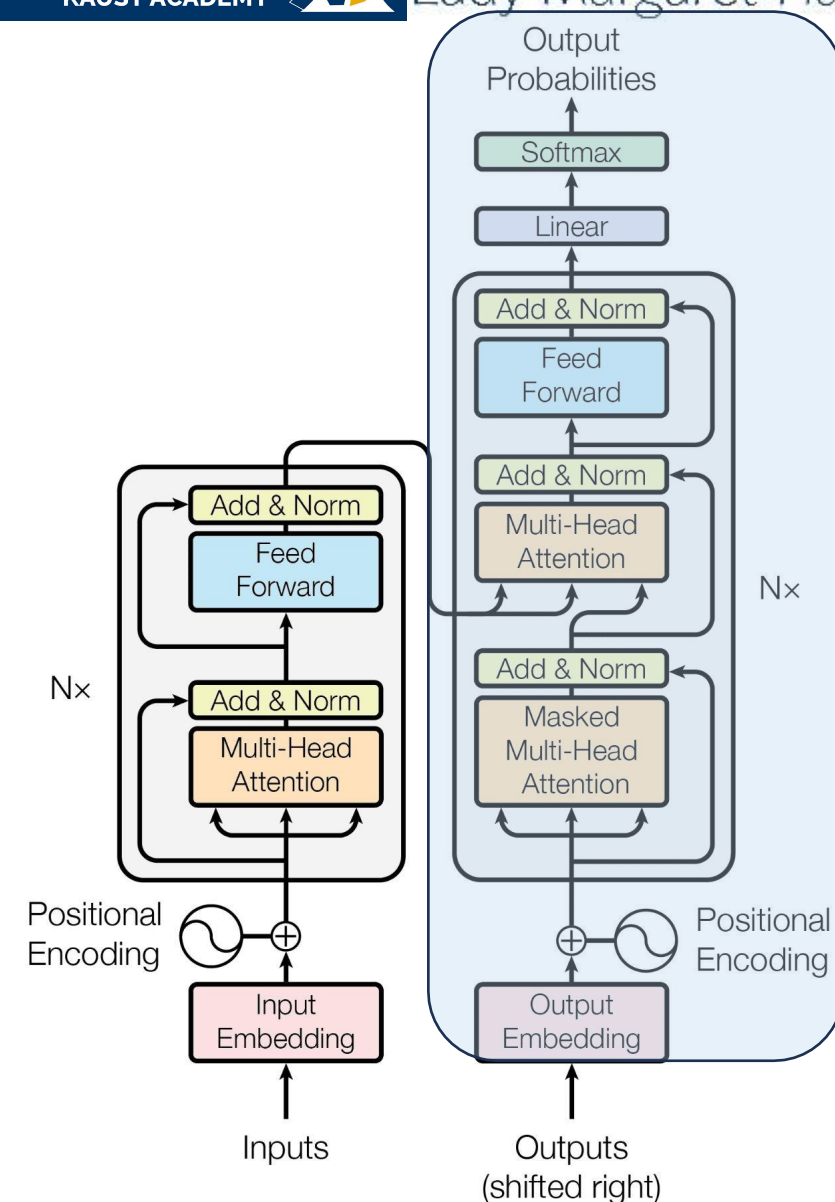
- Similarly, BooksCorpus and English Wikipedia

## GPT Pre-Training Tasks:

- Predict the next token, given the previous tokens
  - More learning signals than MLM

## GPT Pre-Training Results:

- GPT – 117M Params
  - Similarly competitive on GLUE and SQuAD



# GPT – Generative Pretrained Transformer

## GPT Fine-Tuning:

- Prompt-format task-specific text as a continuous stream for the model to fit

### QA

#### Summarization

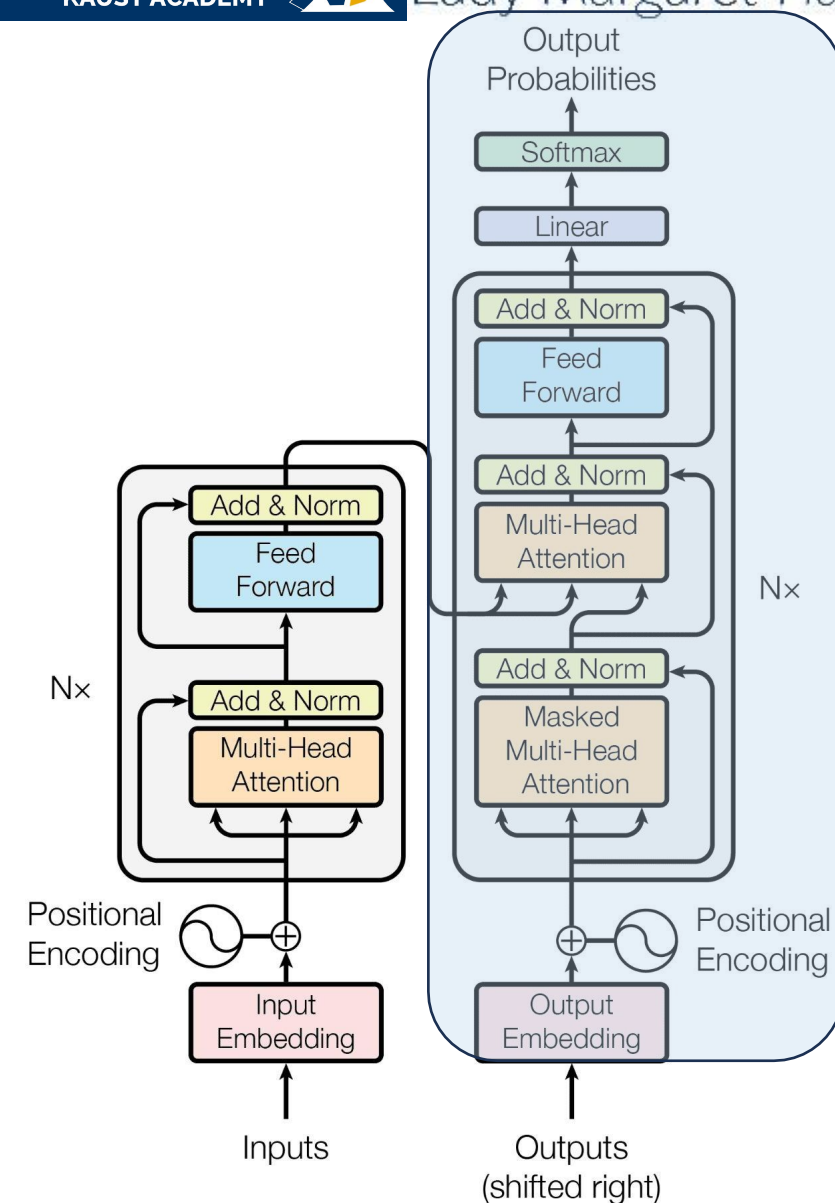
Summarize this article:

Answer the question based on the context.

Context:

Question:

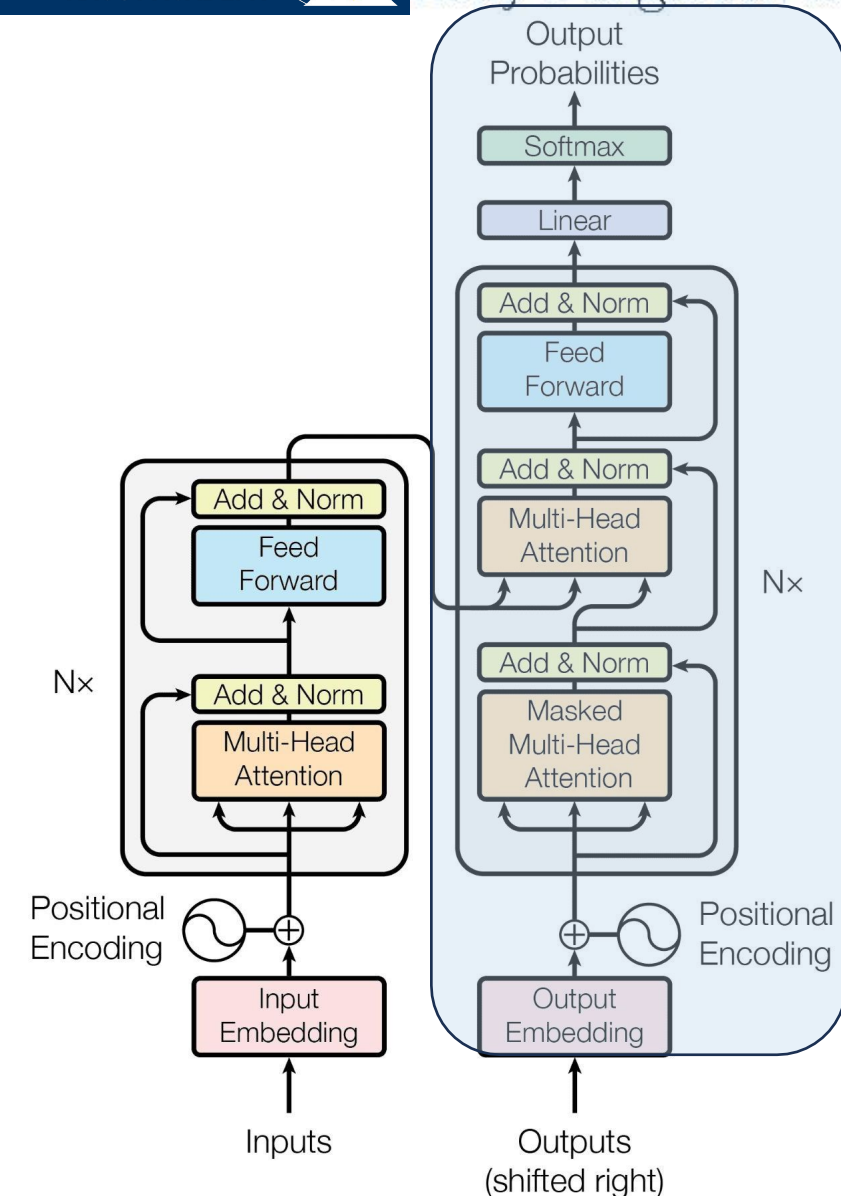
Answer:



# GPT – **Generative** Pretrained Transformer

What is our takeaway from GPT?

- **The Effectiveness of Self-Supervised Learning**
  - Specifically, the model seems to be able to learn from generating the language *itself*, rather than from any specific task we might cook up.






# GPT – **Generative** Pretrained Transformer

What is our takeaway from GPT?


- **The Effectiveness of Self-Supervised Learning**
  - Specifically, the model seems to be able to learn from generating the language *itself*, rather than from any specific task we might cook up.
- **Language Model as a Knowledge Base**
  - Specifically, a generatively pretrained model seems to have a decent zero-shot performance on a range of NLP tasks.



# GPT – **Generative** Pretrained Transformer

## What is our takeaway from GPT?

- **The Effectiveness of Self-Supervised Learning**
  - Specifically, the model seems to be able to learn from generating the language *itself*, rather than from any specific task we might cook up.
- **Language Model as a Knowledge Base**
  - Specifically, a generatively pretrained model seems to have a decent zero-shot performance on a range of NLP tasks.
- **And scaling works!!!**




## ► Key Differences:

Feature	BERT			GPT		
Directionality	Bidirectional			Unidirectional		
Objective	Masked	Language	Modeling	Causal	Language	Modeling
	(MLM)			(CLM)		
Output	Contextual embeddings			Text generation		
Usage	Downstream tasks (e.g., classification, QA)			Generation, few-shot learning		

- ▶ **Kaplan et al. (2020):** “Scaling Laws for Neural Language Models”
- ▶ Performance improves predictably with:
  - More parameters
  - More compute
  - Larger datasets
- ▶ **Optimal allocation of compute:** Train bigger models with less data, rather than small models with lots of data.
- ▶ **Implication:** LLMs like GPT-3 (175B), GPT-4 (est. >500B) are products of scaling laws.

- Scaling improves the perplexity of the LM and improves performance








# Why is this interesting? Look at data scaling



- We know that typical scaling effects look like this when we increase the amount of training data



- Loss and dataset size is linear on a log-log plot
- This is “power-law scaling”




- Can we understand scaling by positing scaling laws ?
- With scaling laws, we can make decisions on architecture, data, hyperparameters by training smaller models
- Open AI Study : **Scaling Laws for Neural Language Models** ([Kaplan et al. 2020](#))

- Open AI Study : **Scaling Laws for Neural Language Models** ([Kaplan et al. 2020](#))
- Key Findings:
  - Performance depends strongly on scale, and weakly on the model shape
  - Larger models are more sample-efficient
  - Smooth power laws ( $y = ax^k$ ) b/w empirical performance & N - parameters, D - dataset size, C - compute

- The effect of some hyperparameters on big LMs can be predicted before training – optimizer (Adam v/s SGD), model depth, LSTM v/s Transformer
- Idea:
  - Train a few smaller models
  - Establish a scaling law (e.g. ADAM vs SGD scaling law)
  - Select optimal hyper param based on the scaling law prediction



# Model Scaling: GPT-3





- Emergent abilities:
  - not present in smaller models but is present in larger models
  - Do LLMs like GPT3 have these ?
- Findings:
  - GPT-3 trained on text can do arithmetic problems like addition and subtraction
  - Different abilities “emerge” at different scales



- Emergent abilities:
  - not present in smaller models but is present in larger models
  - Do LLMs like GPT3 have these ?
- Findings:
  - GPT-3 trained on text can do arithmetic problems like addition and subtraction
  - Different abilities “emerge” at different scales
  - **Model scale is not the only contributor to emergence** – for 14 BIG-Bench tasks, LaMDA 137B and GPT-3 175B models perform at near-random, but PaLM 62B achieves above-random performance
  - Problems LLMs can’t solve today may be emergent for future LLMs

## ► Pre-training Phase:

- Large-scale unsupervised training on corpus (e.g., Common Crawl, Books)
- Objective: learn general-purpose language representations

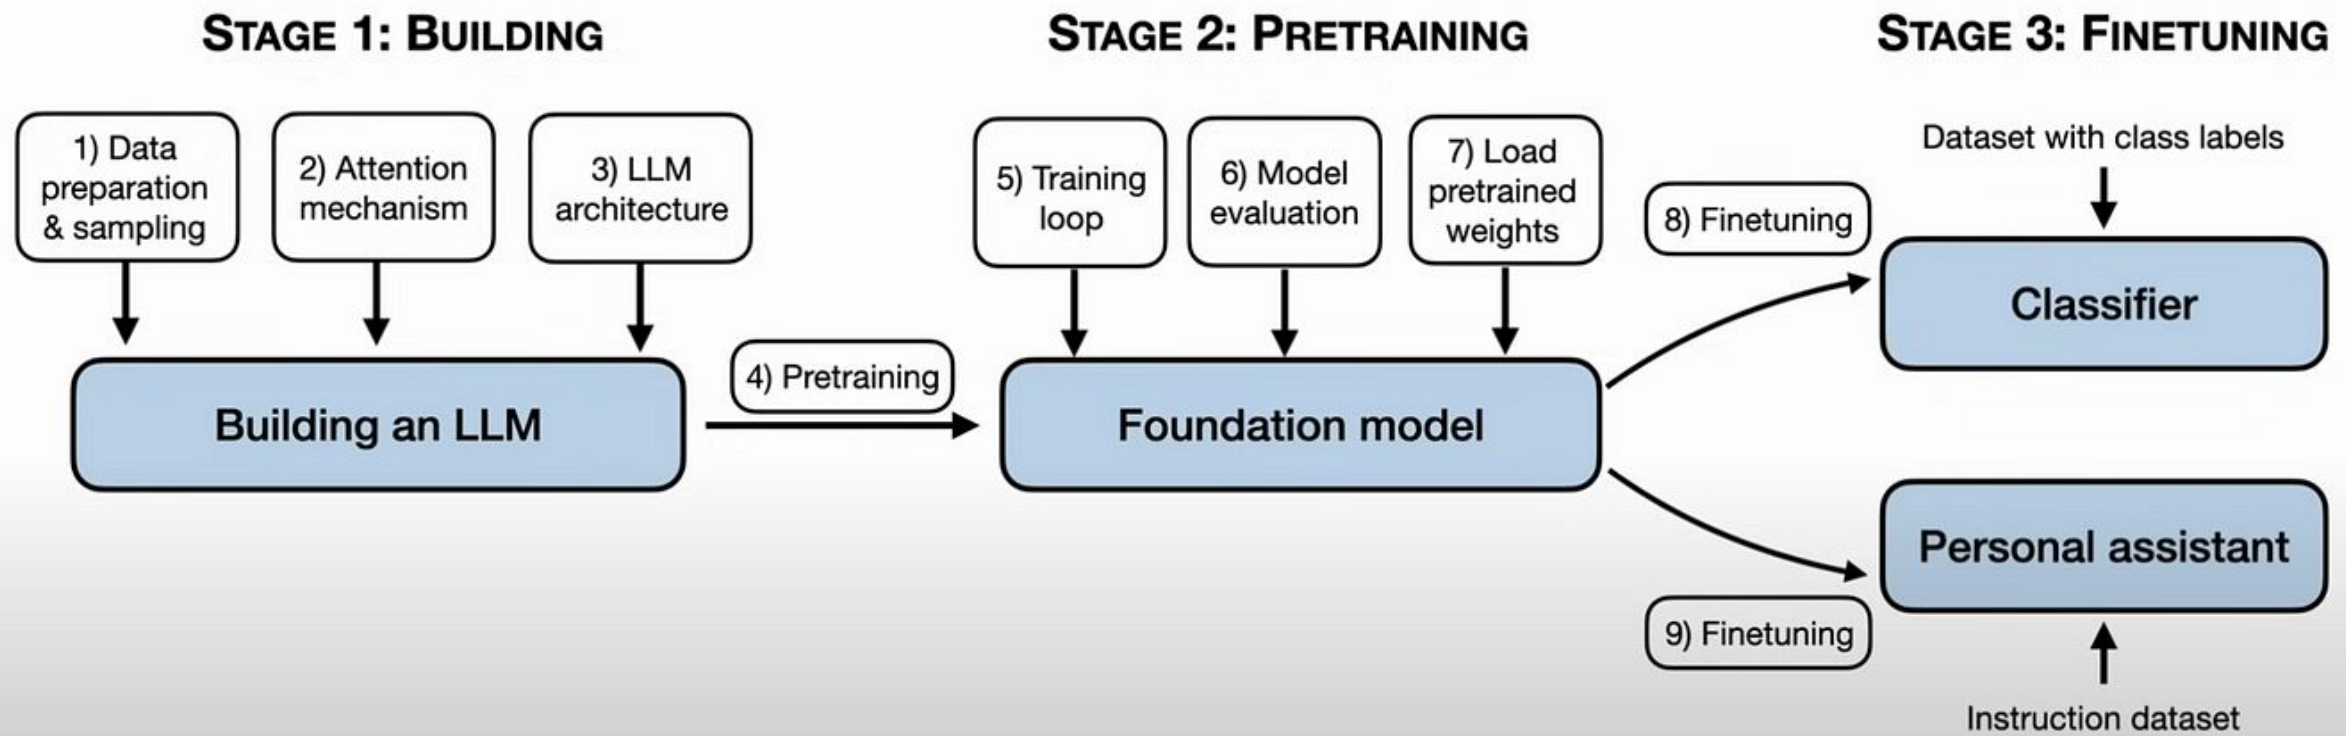
## ► Why Pre-train?

- Data-efficient fine-tuning
- Enables zero-shot and few-shot capabilities
- Foundation for instruction tuning, alignment

## ► Challenges:

- Massive compute costs
- Environmental concerns (carbon footprint)

# Pre-training Overview







1. Auto-regressive Pre-training - Train to predict the next token on very large-scale corpora ( ~3 trillion tokens)

1. Auto-regressive Pre-training - Train to predict the next token on very large scale corpora ( ~3 trillion tokens)
2. Instruction Fine-tuning/ Supervised Fine-tuning (SFT) - Fine-tune the pre-trained model with pairs of (instruction+input,output) with large dataset and then with small high-quality dataset

Instruction fine-tuning provides as a prefix a natural language description of the task along with the input.

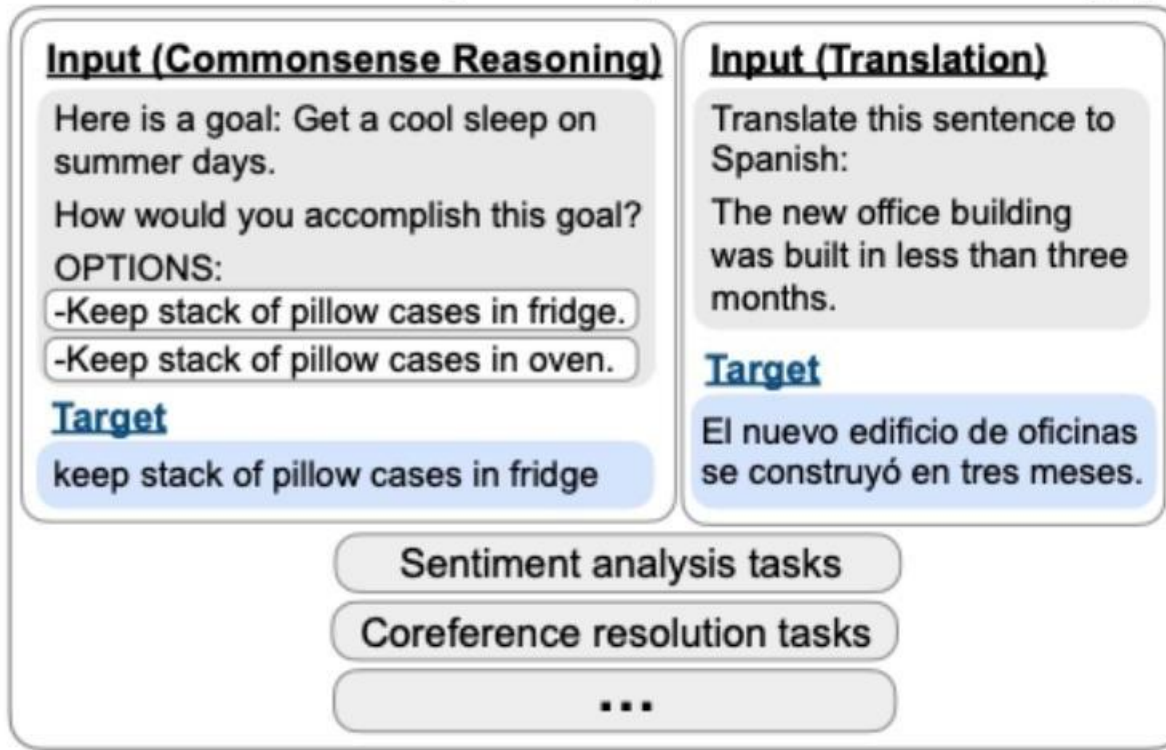
- E.g. Translate into French this sentence: my name is -> je m'appelle



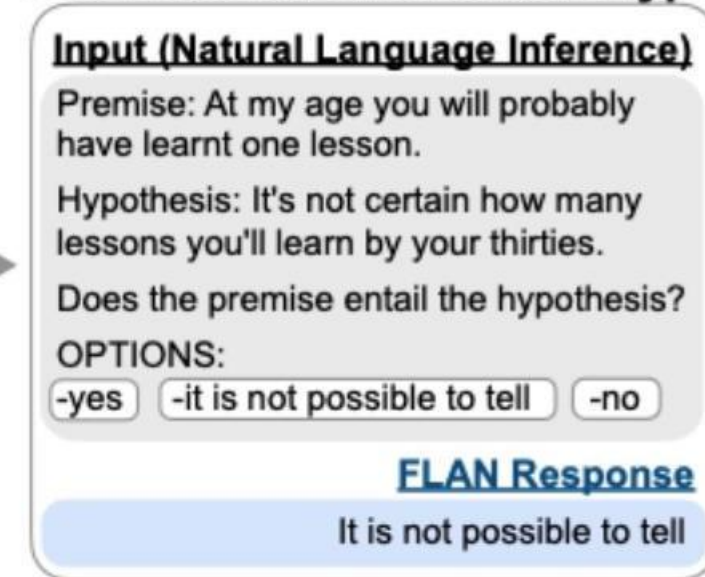
- Objective function
  - Loss computed only for target tokens in SFT, all tokens are targets in pre-training
- Input and Target
  - Instruction + input as input with the target in SFT and only input as input with shifted input as target
- Purpose
  - Pre-training makes good generalist auto-completes but good SFT builds models that can do many unseen tasks
  - SFT can also guide nature of outputs in terms of safety and helpfulness

# Instruction Tuning ([Wei et. al. 2021](#))

## Finetune on many tasks (“instruction-tuning”)



## Inference on unseen task type



- LLMs may produce
  - Harmful text – unparliamentary language, bias and discrimination
  - Text that can cause direct harm – allowing easy access to dangerous information
- Therefore, LLMs should be trained to produce outputs that align with human preferences and values
- Modern LLMs do so by using SFT and by using human preference directly in model training



# Tokenization: Why It Matters?

## ▶ Language $\neq$ Input

- Text must be converted to numbers  $\rightarrow$  tokens.

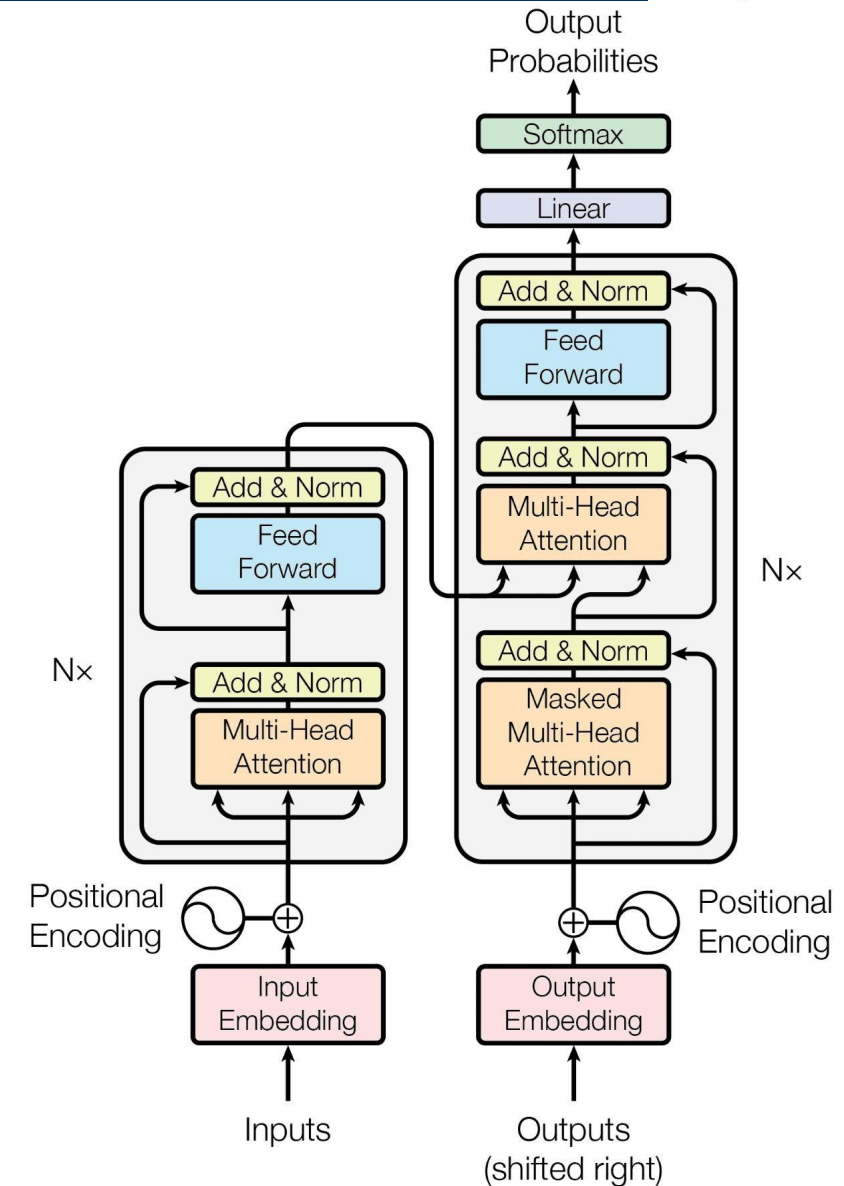
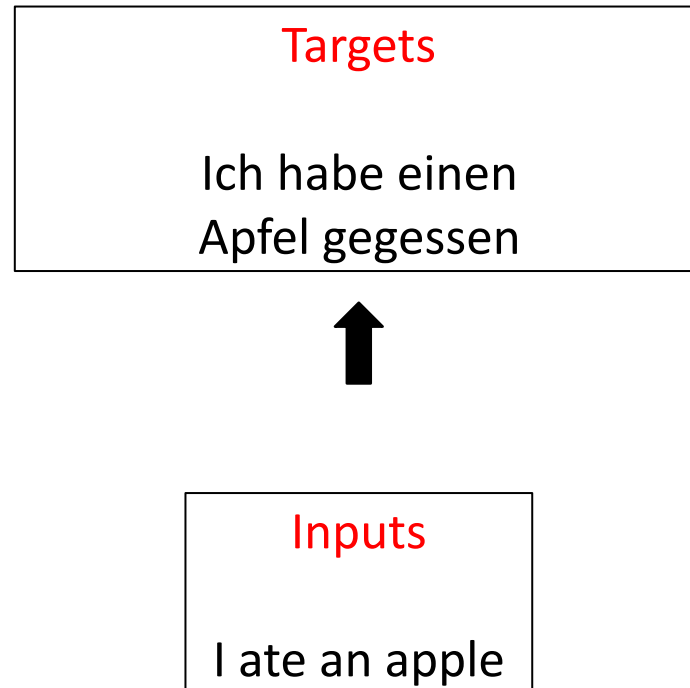
## ▶ Tokenization:

- Splits text into manageable units.
- Balances granularity and vocabulary size.

## ▶ Good tokenizer =

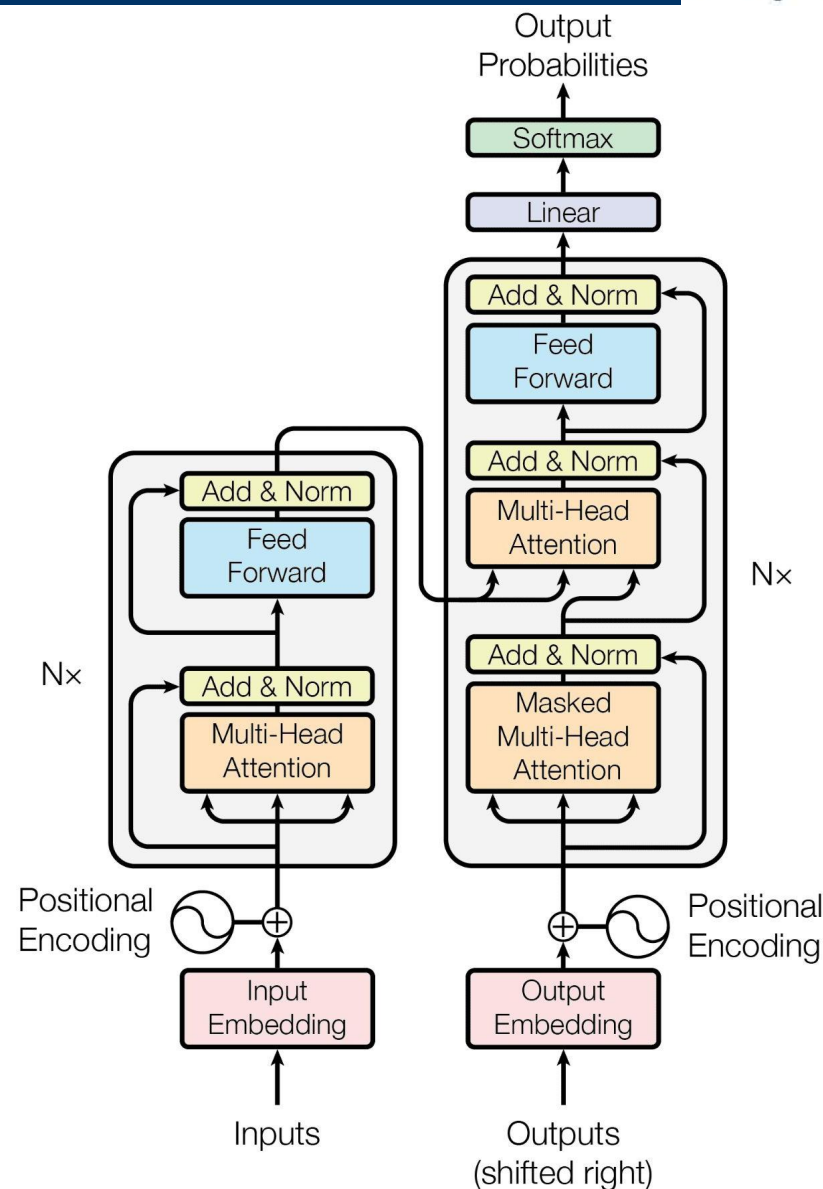
- Efficient sequence length
- High vocabulary coverage
- Robust across domains (code, multilingual, etc.)

# Machine Translation



## Processing Inputs

Inputs  
I ate an apple

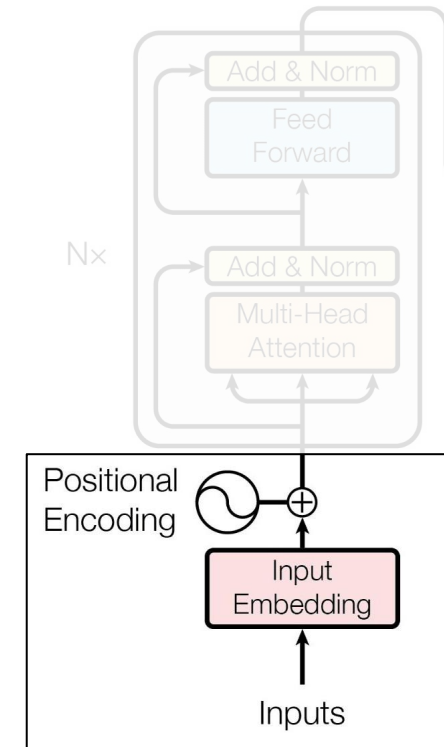


# Tokenization

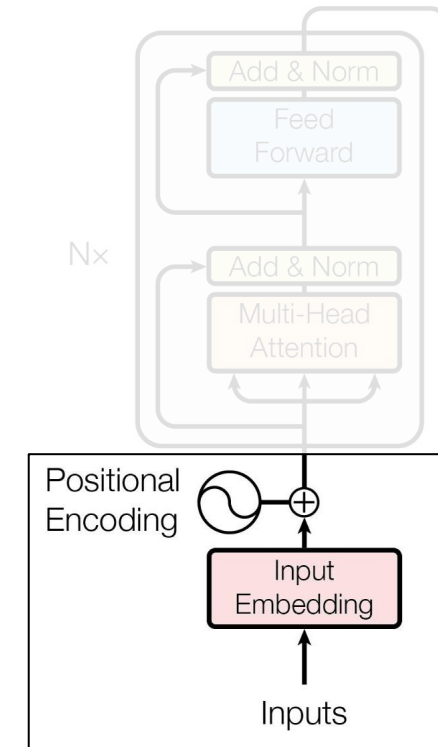
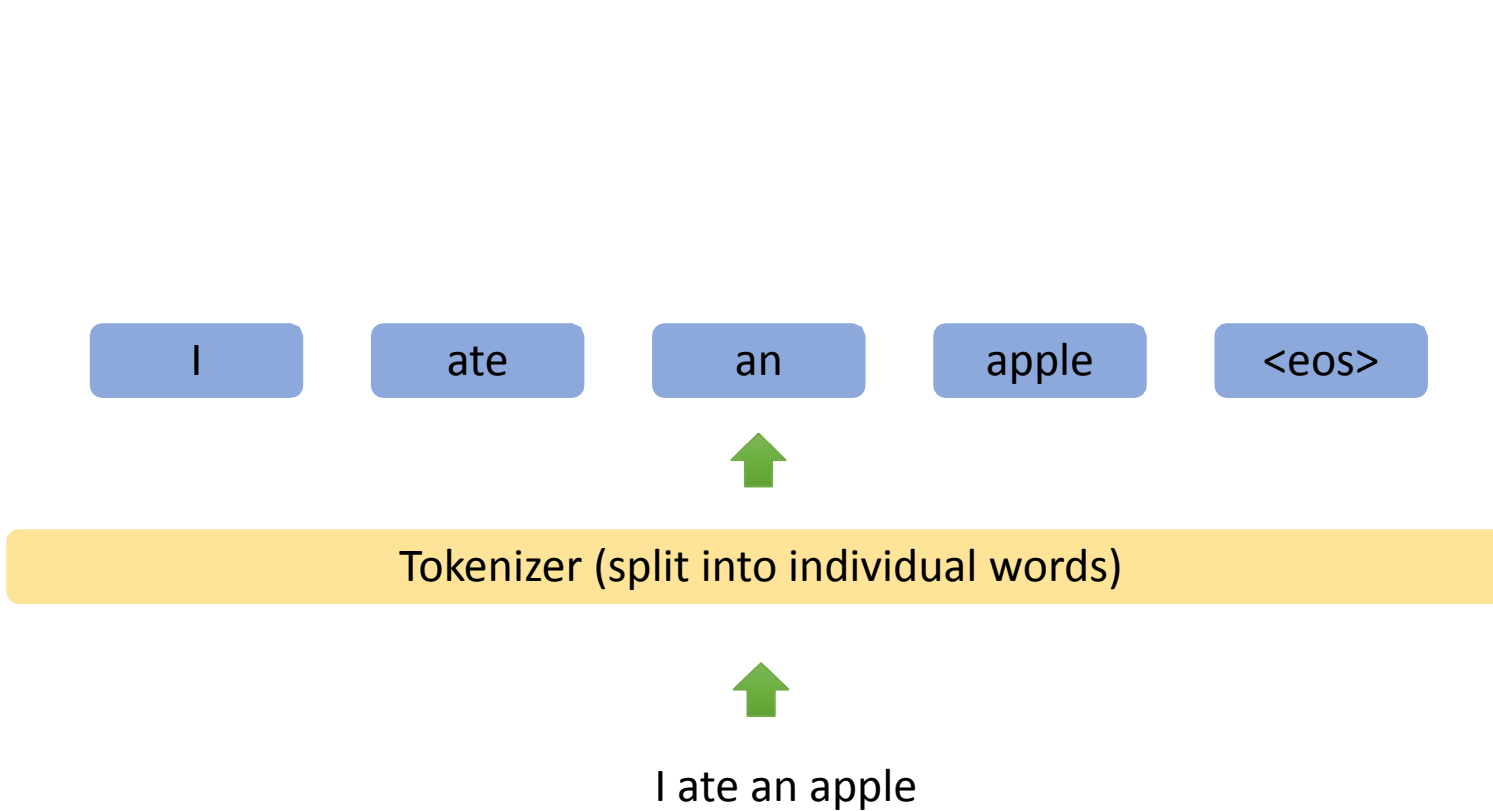
Tokenizer (split into individual words)



I ate an apple



# Tokenization





## ▶ Character-level

- Fine-grained, handles unknowns
- Long sequences, slow training

## ▶ Word-level

- Intuitive, natural boundaries
- Large vocab, OOV (Out-of-Vocab) issues

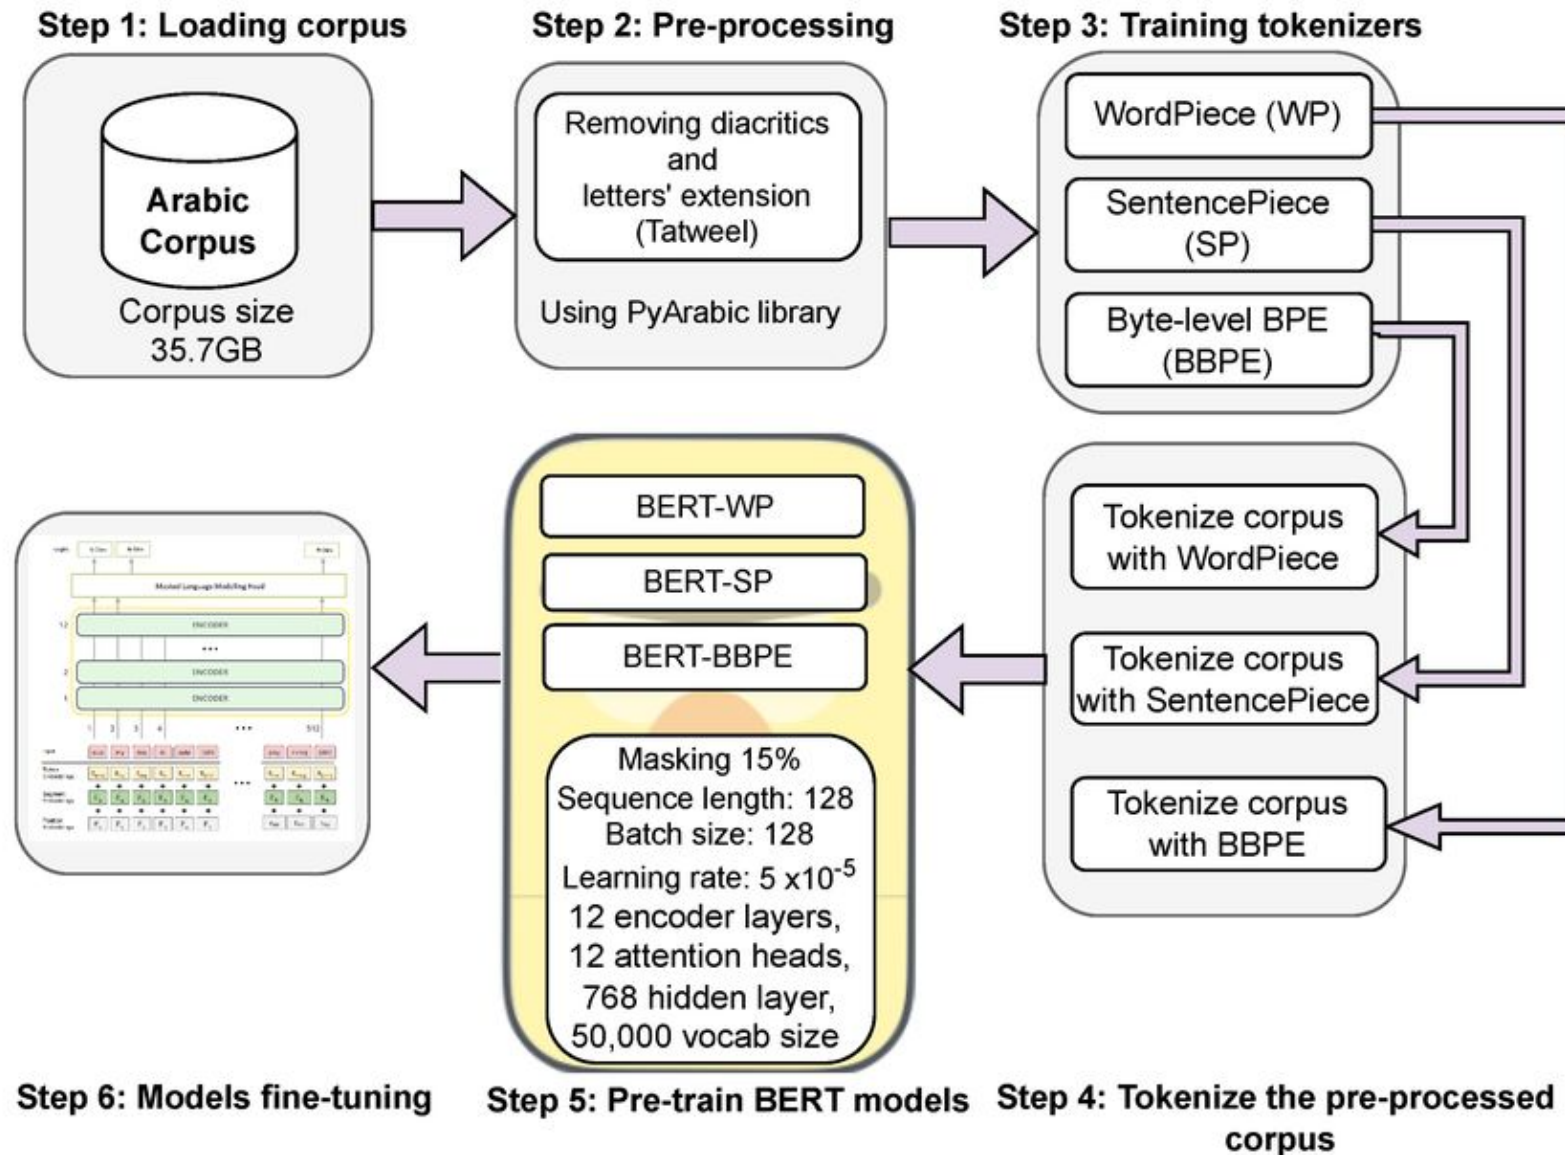
## ▶ Subword-level

- Balance of generalization & compactness
- Requires segmentation algorithm

## ▶ SentencePiece/Unigram LM

- Learned from data, language-agnostic

# Different Levels of Tokenization



# Different Levels of Tokenization

"Machine",  
"learning",  
"is", "fun", "."

**Word  
Tokenization**

"ma",  
"chine",  
"learn", "ing",

**Character  
Tokenization**

"M", "a", "c",  
"h", "i", "n",  
"e", "l", "e",  
"a", "r", "n",  
"i", "n", "g"

**Tokenization Of  
Subwords**

## ► How BPE works:

- Start with characters.
- Merge most frequent pair (e.g., "t", "h" → "th").
- Repeat until vocab size reached.

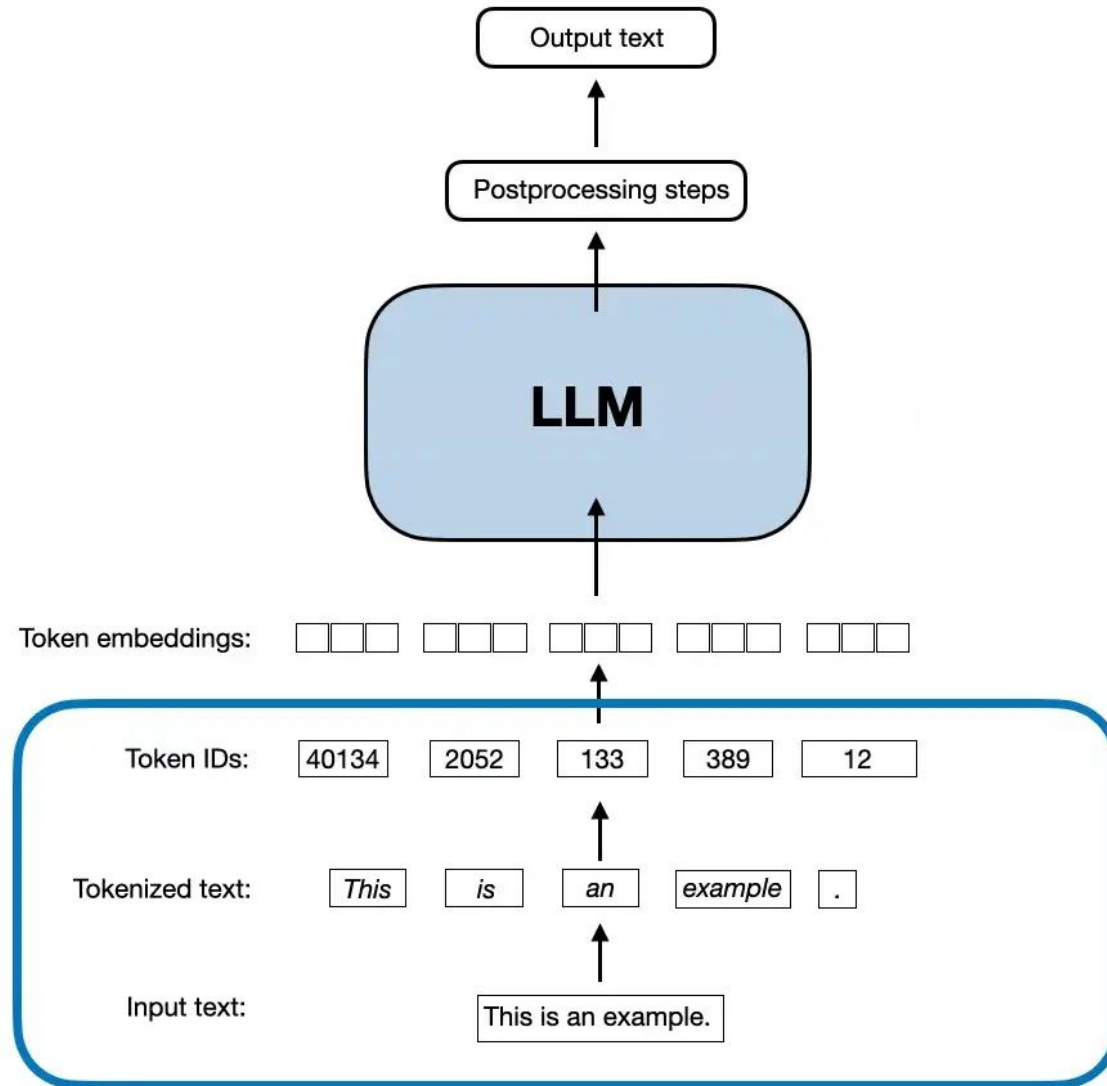
## ► Advantages:

- Handles unknown words well (e.g., unbelievably → un + believ + ably)
- Reduces sequence length

## ► Limitations:

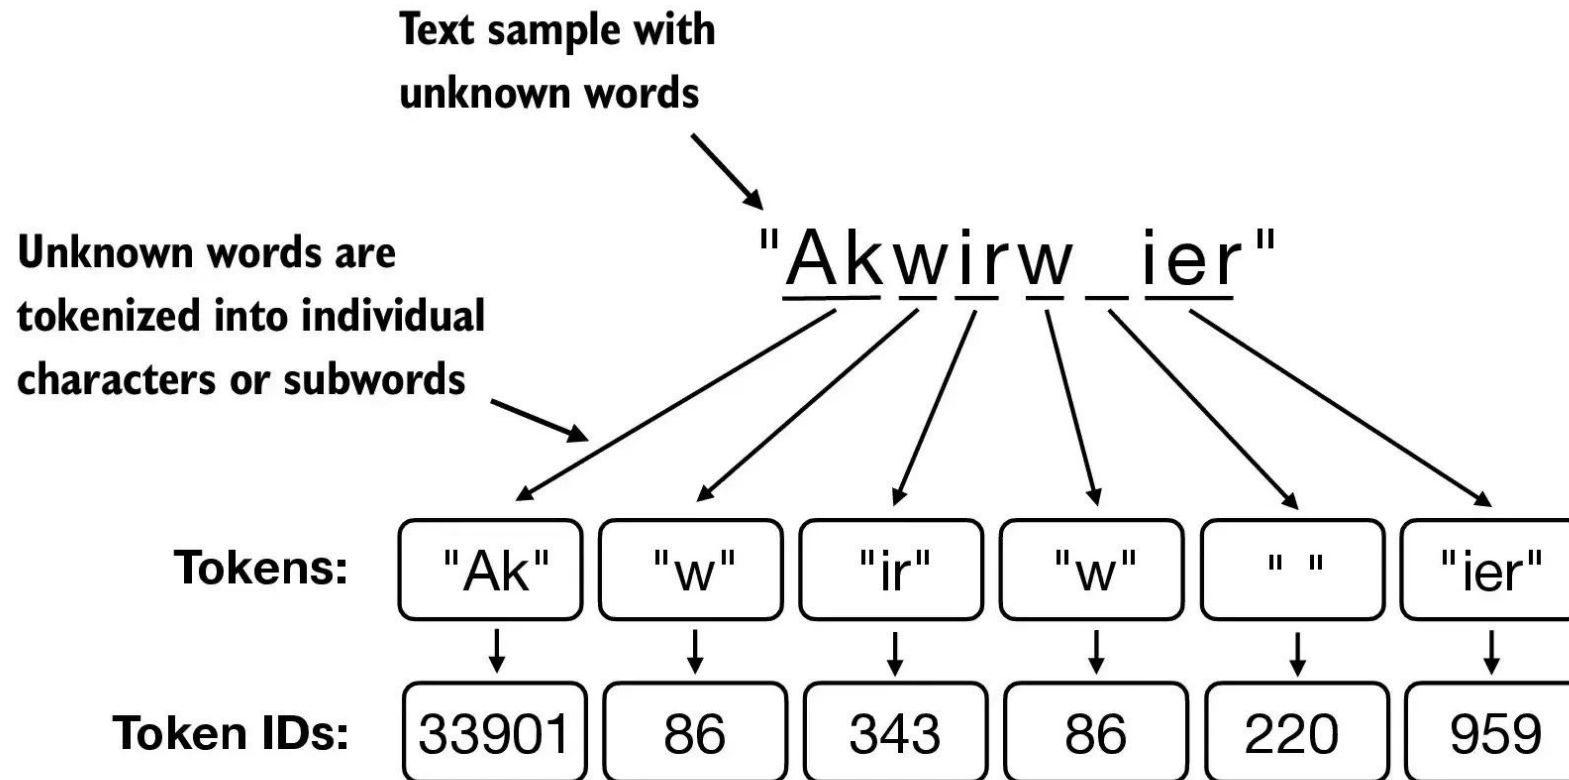
- Not optimal for all scripts (e.g., Chinese)
- Word boundaries may be unclear

# Byte-Pair Encoding (BPE)





# Byte-Pair Encoding (BPE)





# Byte-Pair Encoding (BPE)

## How BPE Works

Start with Characters

Count Frequent Pairs

Merge Pairs

Repeat Process

## Why Use BPE

Efficiency

Generalization

Byte Pair  
Encoding  
(BPE)

## Example

Text: "low lower"

Split: ['l', 'o', 'w'], ['l', 'o', 'w', 'e', 'r']

Merge: ('l', 'o') → ['lo', 'w'], ['lo', 'w', 'e', 'r']

Result: Tokens → ['lo', 'w', 'lo', 'w', 'e', 'r']

## Applications

Tokenization in LLMs

Handling Rare Words

## ► What are Token-Free LLMs?

- Models that operate directly on raw text without tokenization.
- Aim to eliminate the need for pre-defined tokens.

## ► Advantages:

- Avoids tokenization errors and biases.
- Can handle arbitrary text lengths and formats.
- Potentially more efficient for certain tasks.

## ► Challenges:

- Requires novel architectures to process raw text effectively.
- May struggle with long-range dependencies without tokenization.

## ► Example: ByT5 (Xue et al., 2022)

- Character-level variant of T5 that processes raw bytes instead of tokens.
- Eliminates the need for a tokenizer, enabling better multilingual and low-resource language support.

## ► Pros:

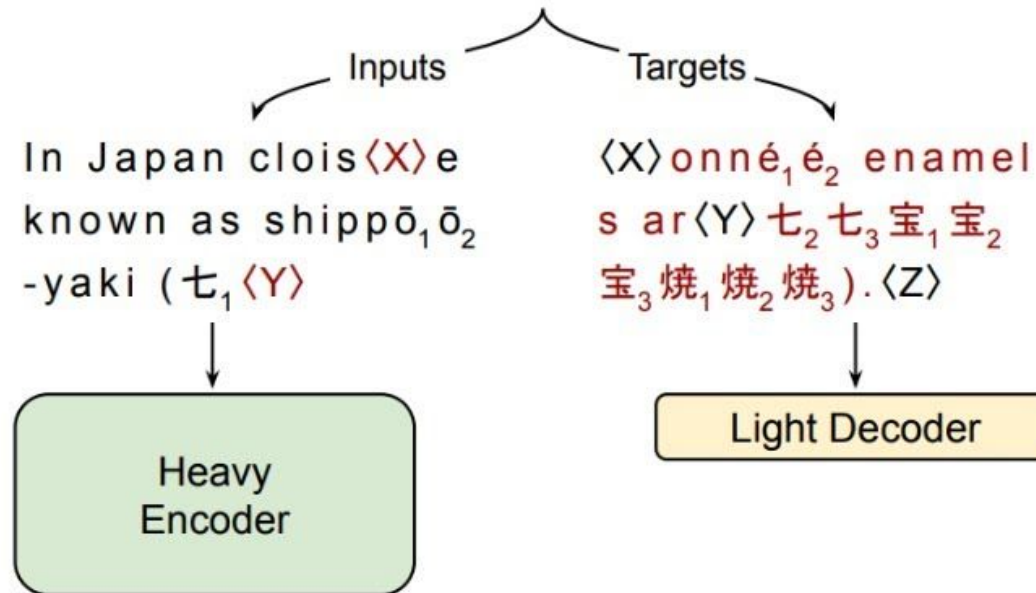
- No preprocessing or tokenization pipeline required.
- Handles any language or script without modification.
- Performs better on low-resource and unseen languages.

## ► Cons:

- Training is slower due to longer input sequences.
- Increased computational requirements.
- May require more data to achieve comparable performance.

## ByT5: Towards a Token-Free Future with Pre-trained Byte-to-Byte Models

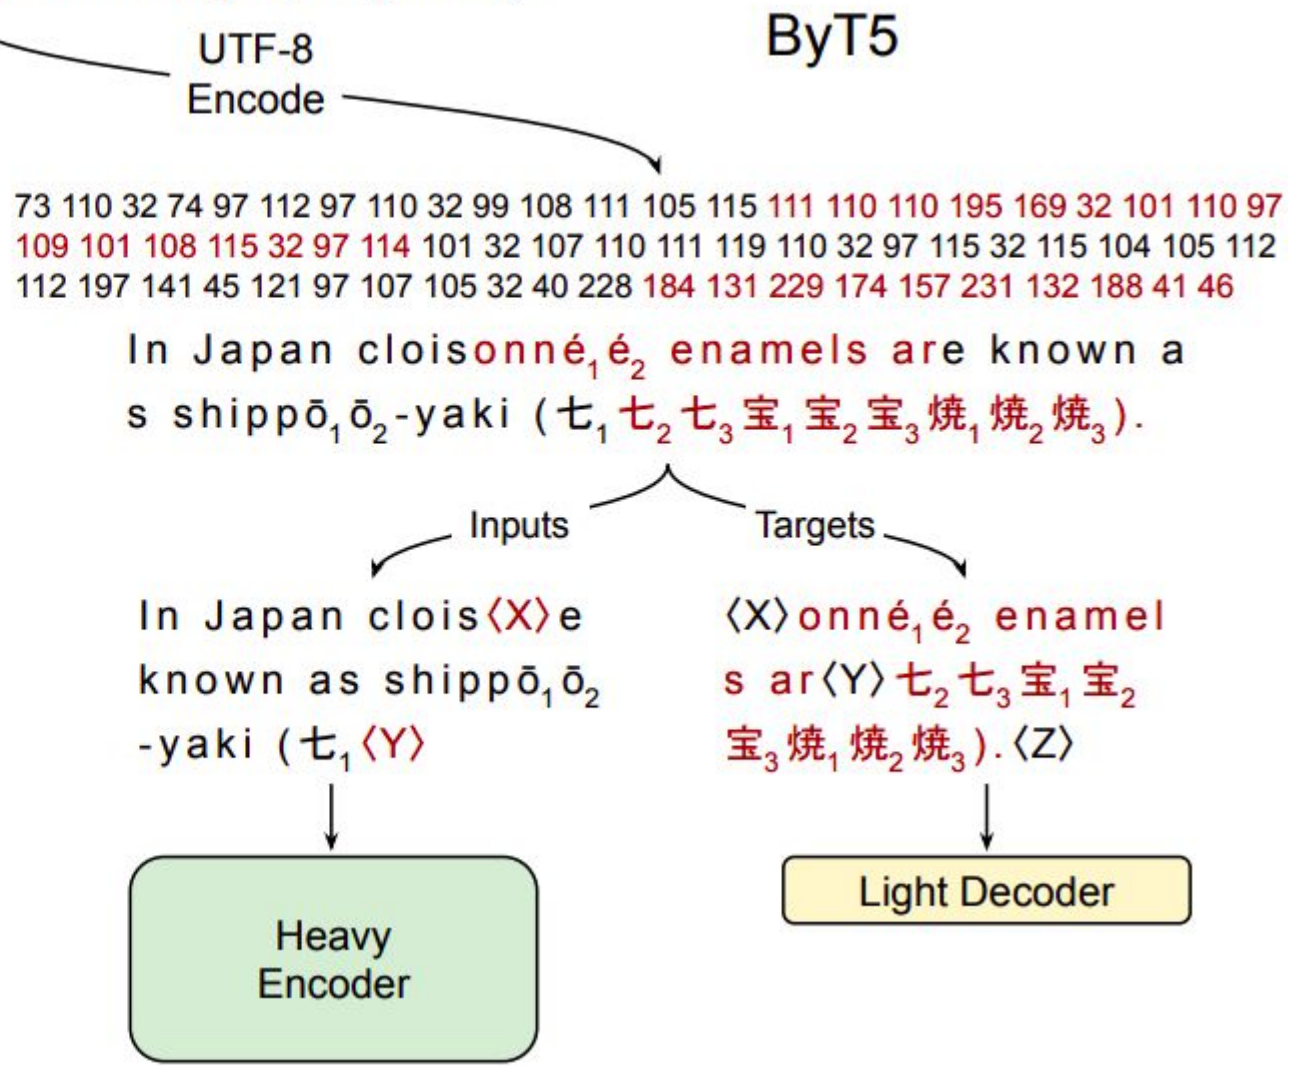
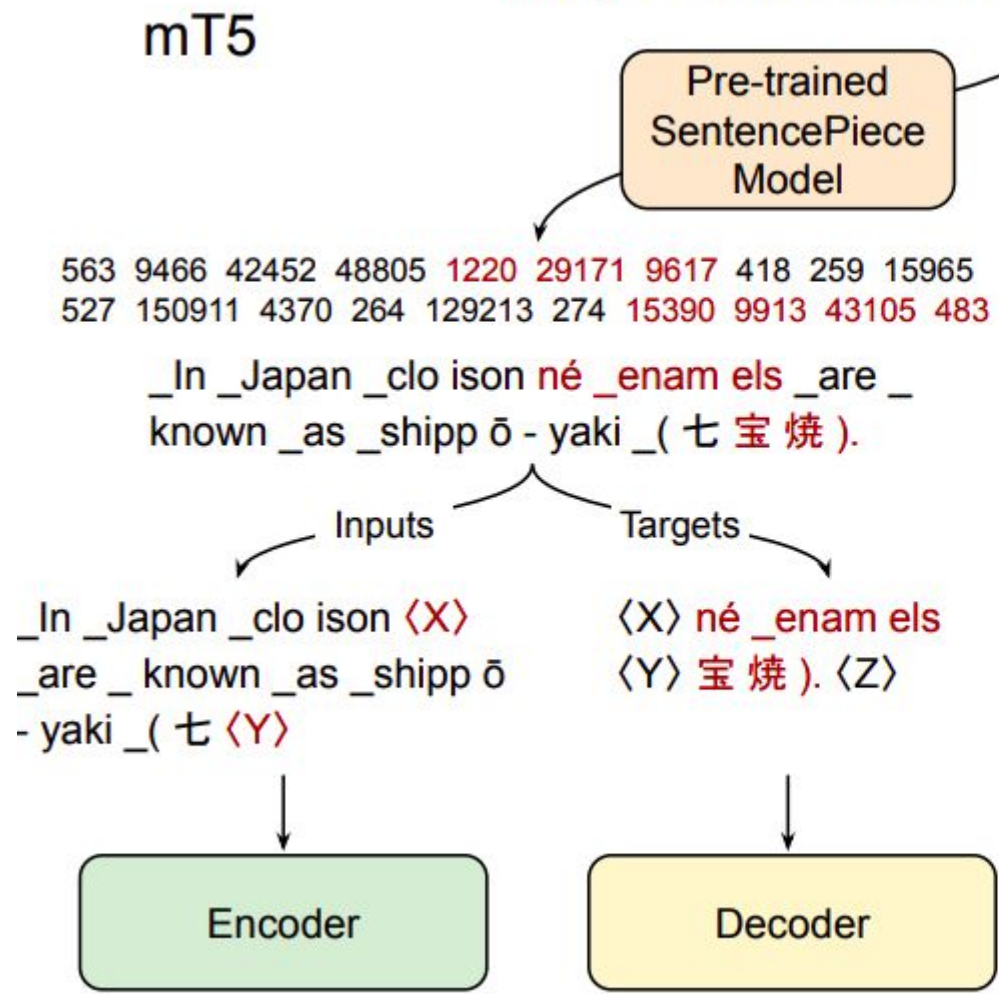
s shippō<sub>1</sub>ō<sub>2</sub>-yaki (七<sub>1</sub>七<sub>2</sub>七<sub>3</sub>宝<sub>1</sub>宝<sub>2</sub>宝<sub>3</sub>焼<sub>1</sub>焼<sub>2</sub>焼<sub>3</sub>).



# NLP Journal Club



In Japan cloisonné enamels are known as shippō-yaki (七宝焼).






Figure 3: Per-language performance gaps between ByT5-Large and mT5-Large, as a function of each language’s “compression rate”. **Top:** TyDiQA-GoldP gap. **Bottom:** XNLI zero-shot gap.



## ► Current Research:

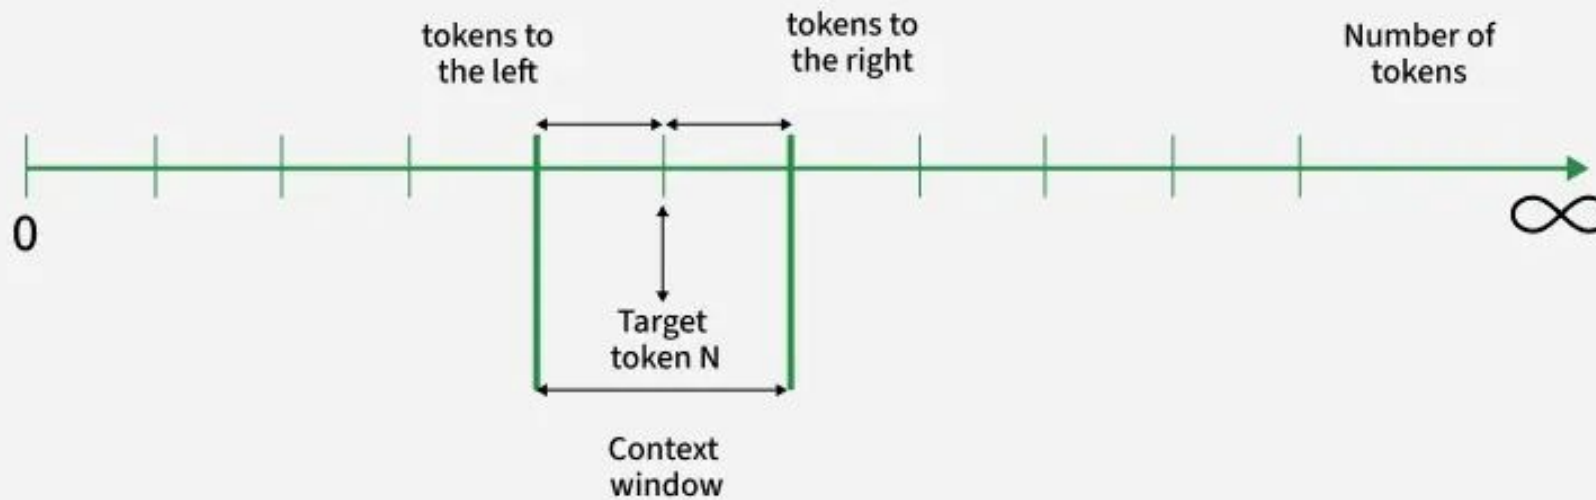
- Exploring architectures like *Raw Transformer* and *Text2Vec*.
- Investigating how to maintain performance without tokenization.

## ► Future Directions:

- Integrating token-free approaches with existing LLMs.
- Enhancing efficiency and scalability of token-free models.

- ▶ Current LLMs have a finite “context window”:
  - GPT-3: 2K tokens
  - GPT-4: Up to 128K tokens
  - Claude 3.5: Up to 200K tokens
- ▶ **Problems:**
  - Long documents get truncated
  - Token limit affects reasoning
  - Difficult to do document-level QA or summarization

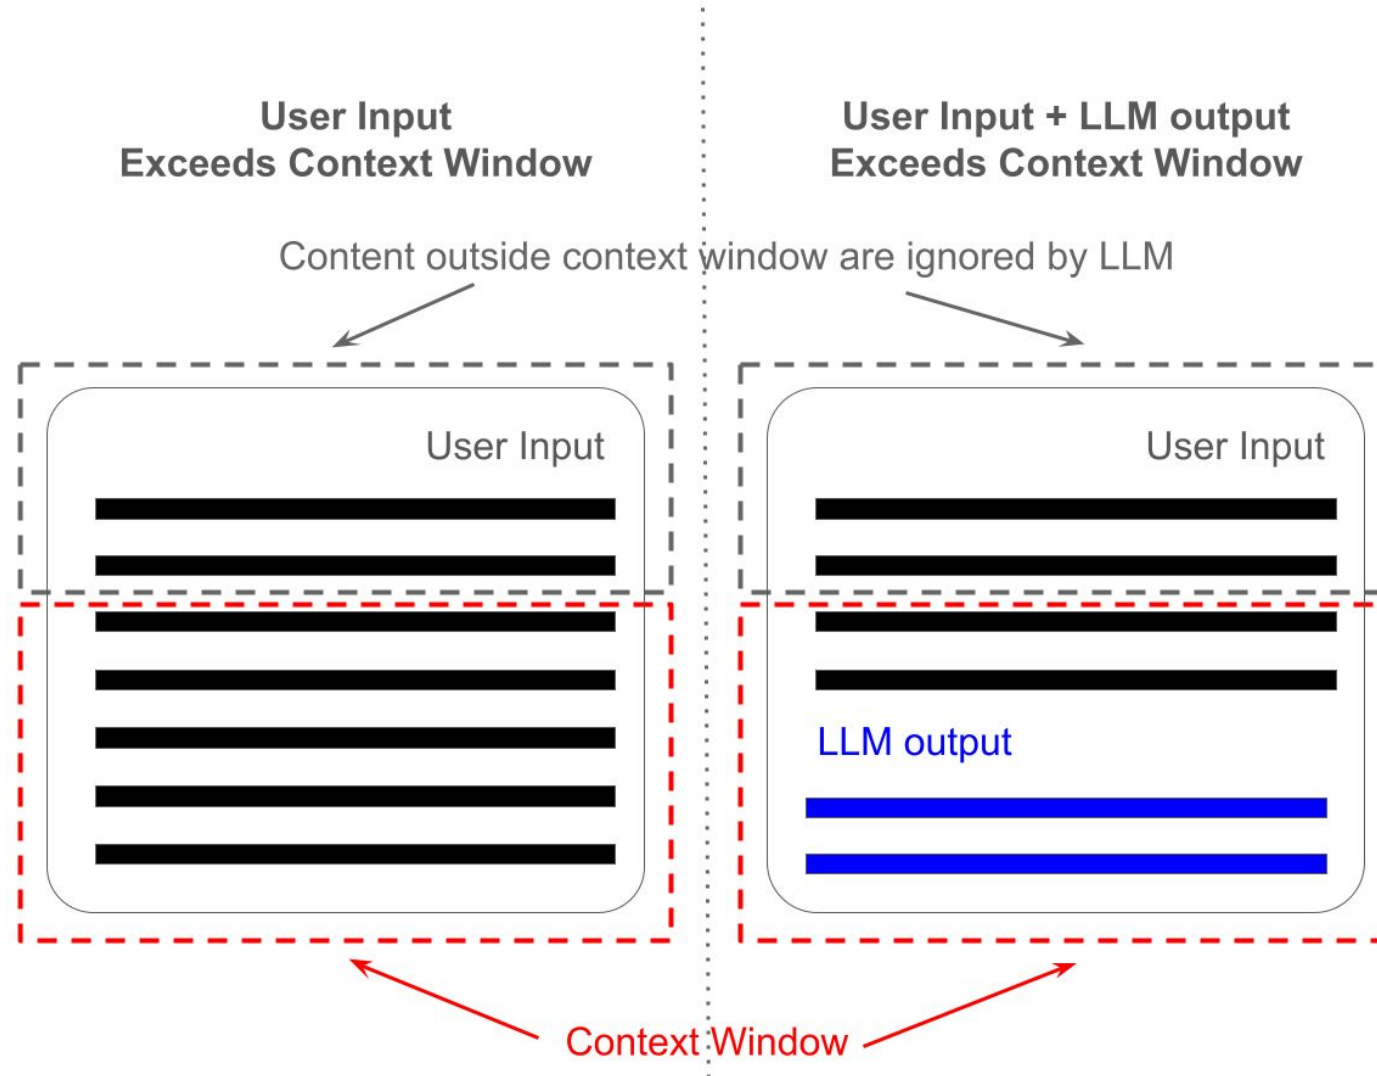
## Example of a Context Window



## Why **Bigger Context Windows** Aren't Always Better?



# Context Window Problems







## ▶ **Efficient Attention Variants:**

- Longformer, BigBird, FlashAttention-2

## ▶ **Memory-Augmented Models:**

- Retrieval-augmented generation (RAG)
- Memory layers (e.g., RetNet)

## ▶ **Chunking + Re-ranking:**

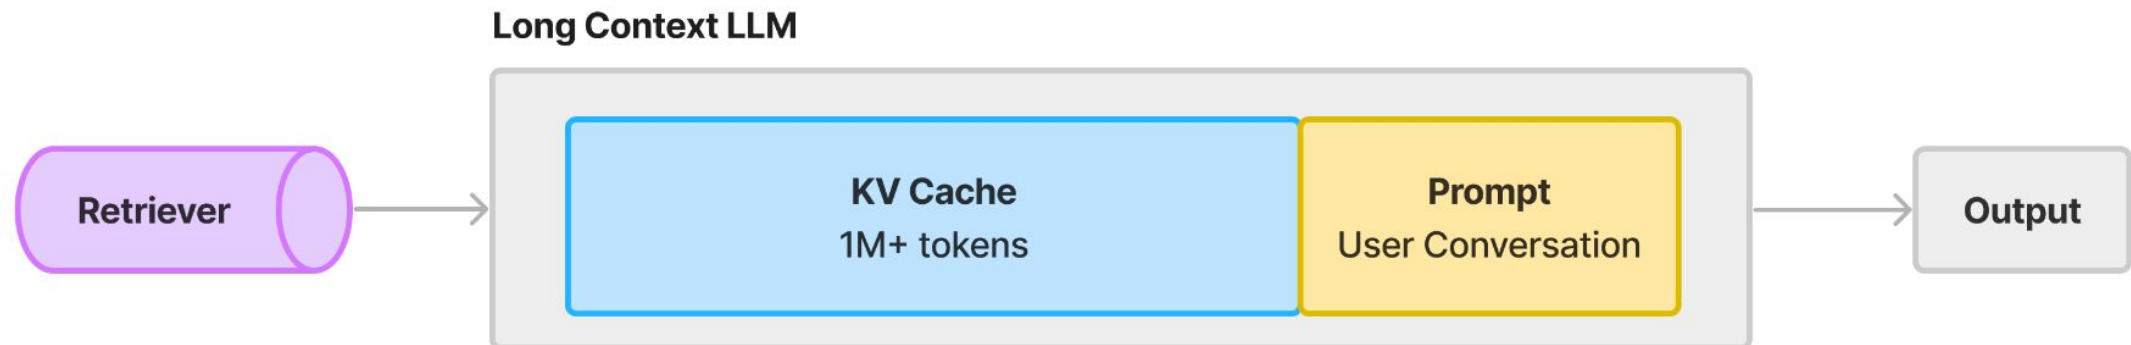
- Process in parts, summarize/join results

## ▶ **Recurrence or State Passing:**

- Transformer-XL, RWKV (RNN-inspired)



# Moving Towards Larger Context Windows



- ▶ High inference costs & latency
- ▶ Bias and toxic outputs
- ▶ Hallucination and factual inconsistency
- ▶ Limited interpretability
- ▶ Require vast pre-training data
- ▶ Poor domain generalization (medical, legal, etc.)

- ▶ **Long Context Handling:** Efficient token reuse, sparse attention, recurrence.
- ▶ **Token-Free Models:** Improved byte/character models.
- ▶ **Multimodal LLMs:** Integrating vision, speech, and more.
- ▶ **Smaller, Efficient LLMs:** Distillation, quantization, sparse models.
- ▶ **Open-Weight & Ethical LLMs:** Responsible, accessible models.

- ▶ BERT and GPT are foundational LLM architectures.
- ▶ Scaling laws dictate optimal growth paths.
- ▶ Tokenization is central to model performance.
- ▶ Context window size remains a bottleneck.
- ▶ Innovations in token-free modeling and memory-efficient transformers are shaping the future.

These slides have been adapted from

- Bhiksha Raj & Rita Singh, 11-785 Introduction to Deep Learning, CMU



- [1] Kaplan et al., “Scaling Laws for Neural Language Models”, 2020.
- [2] Devlin et al., “BERT: Pre-training of Deep Bidirectional Transformers”, 2018.
- [3] Brown et al., “Language Models are Few-Shot Learners (GPT-3)”, 2020.
- [4] Xue et al., “ByT5: Towards a token-free future with pre-trained byte-to-byte models”, 2022.
- [5] Beltagy et al., “Longformer: The Long-Document Transformer”, 2020.
- [6] Tay et al., “Efficient Transformers: A Survey”, 2020.
- [7] OpenAI, “Technical Report on GPT-4”, 2023.



- [8] Google Research Blog: Scaling Transformer Models.
- [9] Anthropic, “Claude 3.5 Release Notes”, 2024.
- [10] FlashAttention: <https://arxiv.org/abs/2205.14135>

## Credits

Dr. Prashant Aparajeya

Computer Vision Scientist — Director(AISimply Ltd)

[p.aparajeya@aisimply.uk](mailto:p.aparajeya@aisimply.uk)

This project benefited from external collaboration, and we acknowledge their contribution with gratitude.